


# Онлайн образование

[otus.ru](https://otus.ru)



**Проверить, идет ли запись**

**Меня хорошо видно  
&& слышно?**





# Couchbase

**Аристов Евгений**

telegram @AEugene

<https://aristov.tech>

# Правила вебинара



Активно участвуем



Задаем вопрос в чат



Вопросы вижу в чате, могу ответить не сразу



# Маршрут вебинара



# Цели вебинара

- 1 Понять зачем используется Couchbase**
- 2 Познакомиться с архитектурой Couchbase**

# Смысл | Зачем вам это уметь, в результате:


**1 Выбрать оптимальный вариант NoSQL решения**

**2 Уметь базово настроить Couchbase**





# The Modern Database for Enterprise Applications





The foundation of all these apps is a Couchbase database



# A Proven Enterprise Solution Chosen by Industry Leaders



## Retail & E-Commerce



**3 of the Top 103**  
eCommerce  
Companies



## Travel & Hospitality



**of the Top 3**  
GDS  
Companies



## Financial Services



**3 of the Top 3**  
Credit Reporting  
Companies



## Healthcare



**3 Fortune 500**  
Healthcare  
Companies



## Telecom



**6 of the Top 10**  
Broadcast  
Companies



## Media & Entertainment



## Gaming



**6 of the Top 10**  
Online Casino  
Gaming Companies



## Industrial IoT



**2 of the Top 2**  
IoT  
Platforms

# Couchbase Behind Today's Business-Critical Applications



Customers

Application

Performance

**Linked**in

Caching & session store for single view

**2M+**

reads/sec.

**10M**

queries/sec.

**TESCO**

Real-time pricing, product catalog, inventory management

**10M+**

unique SKUs

**35K**

requests/sec.

**amadeus**

Flight availability, booking, pricing analytics, etc.

**15M**

ops / second

**<2.5ms**

response time

**COMCAST**

Customer 360 single view, unified notes

**210M**

documents

**100K**

users

**UNITED**

Real-time crew management, scheduling and resources

**41K**

pilots and crew

**148M**

travelers in 2017

Infrastructure

Developer Agility


Performance at Scale

Manageability

Security + Availability

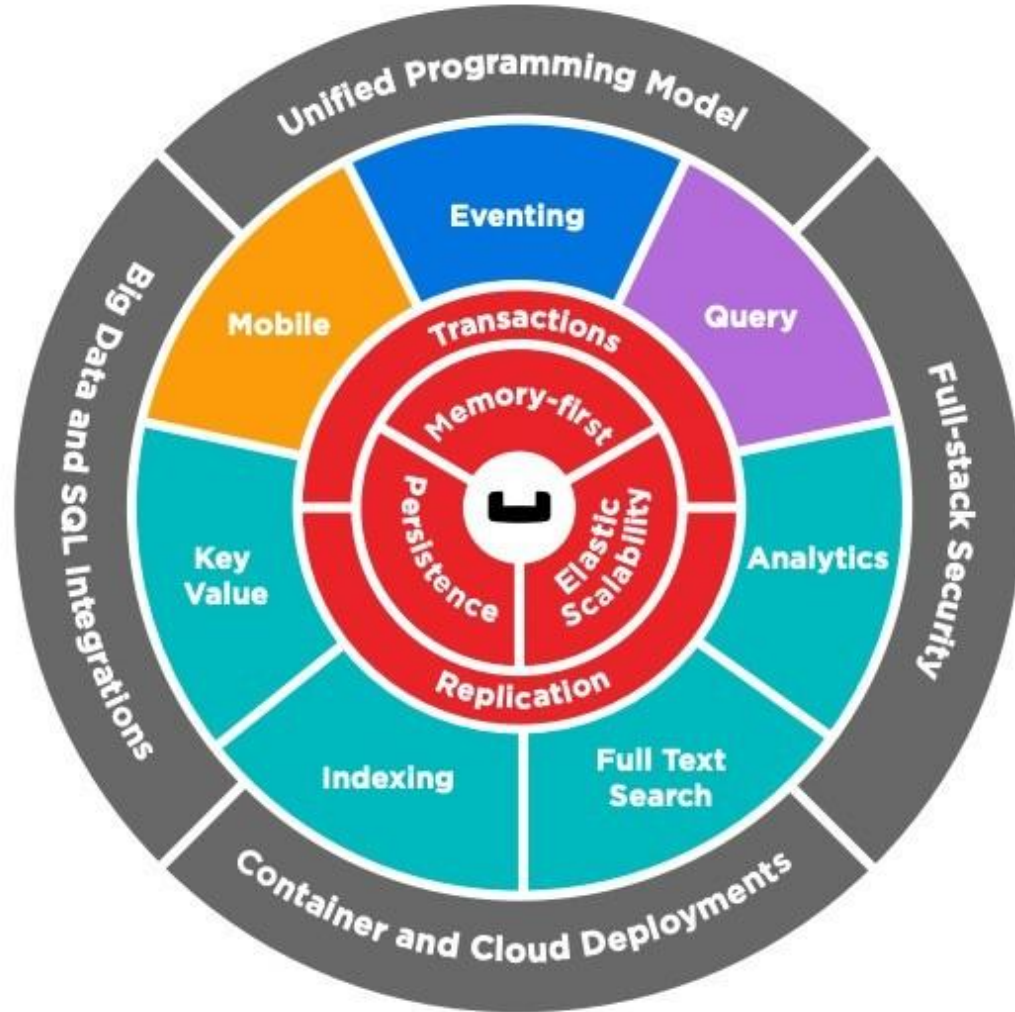


# Couchbase: Brief Company Overview




- Founded 2011
- ~ 700 employees
- 70%+ new business growth
- 100%+ growth in total contract value
- 55%+ billings growth
- 50%+ growth in average deal size
- IPO 22.07.2021

# Architecture Overview



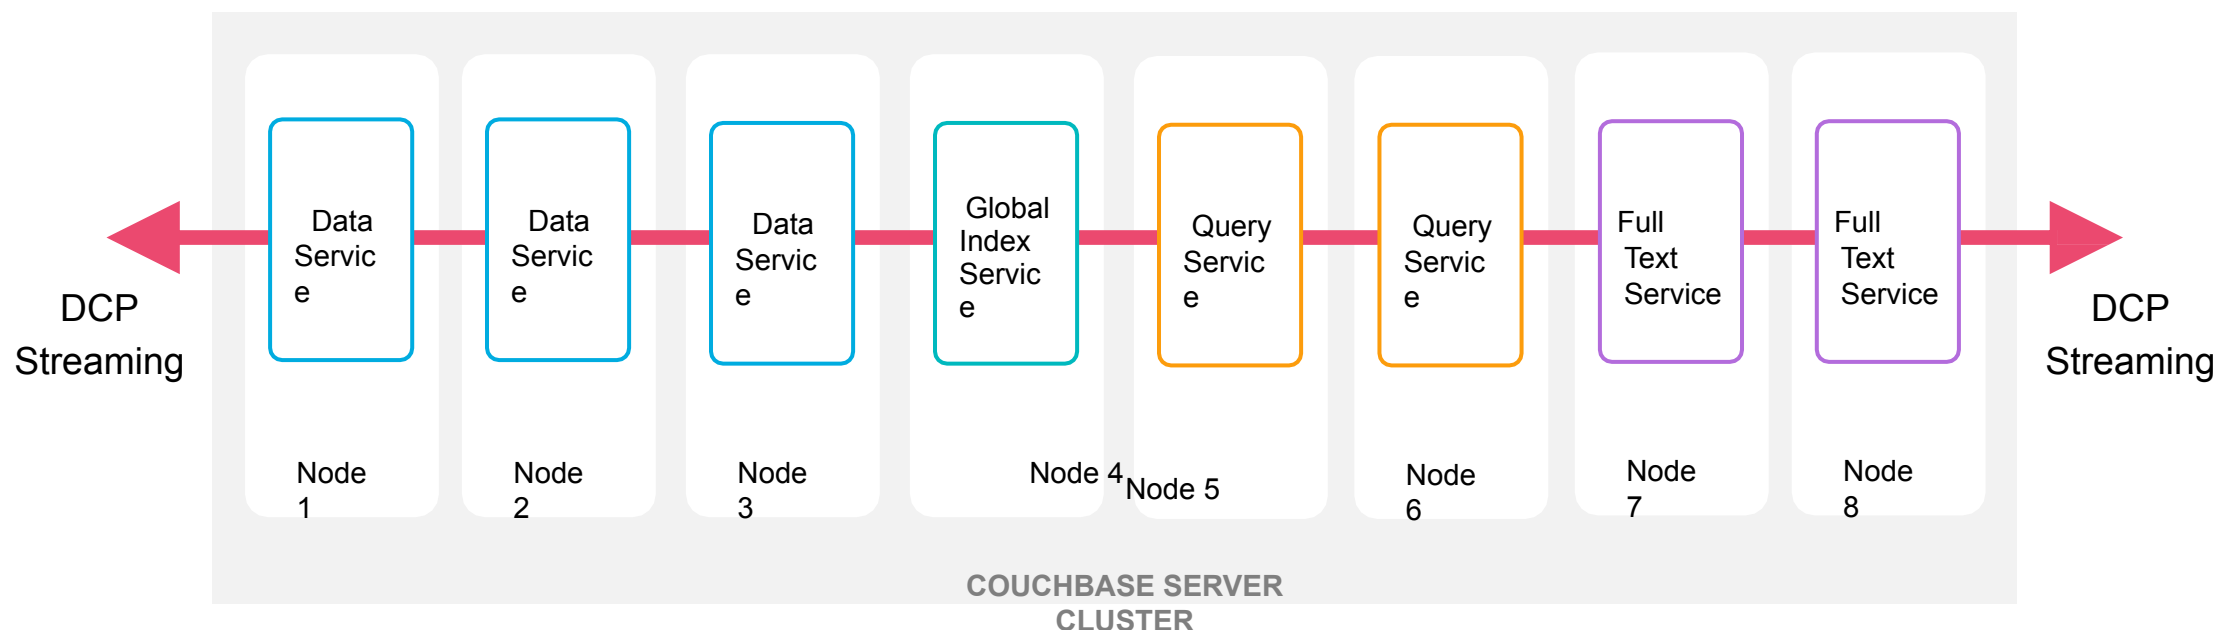
# Memory-First



# Memory-First Architecture

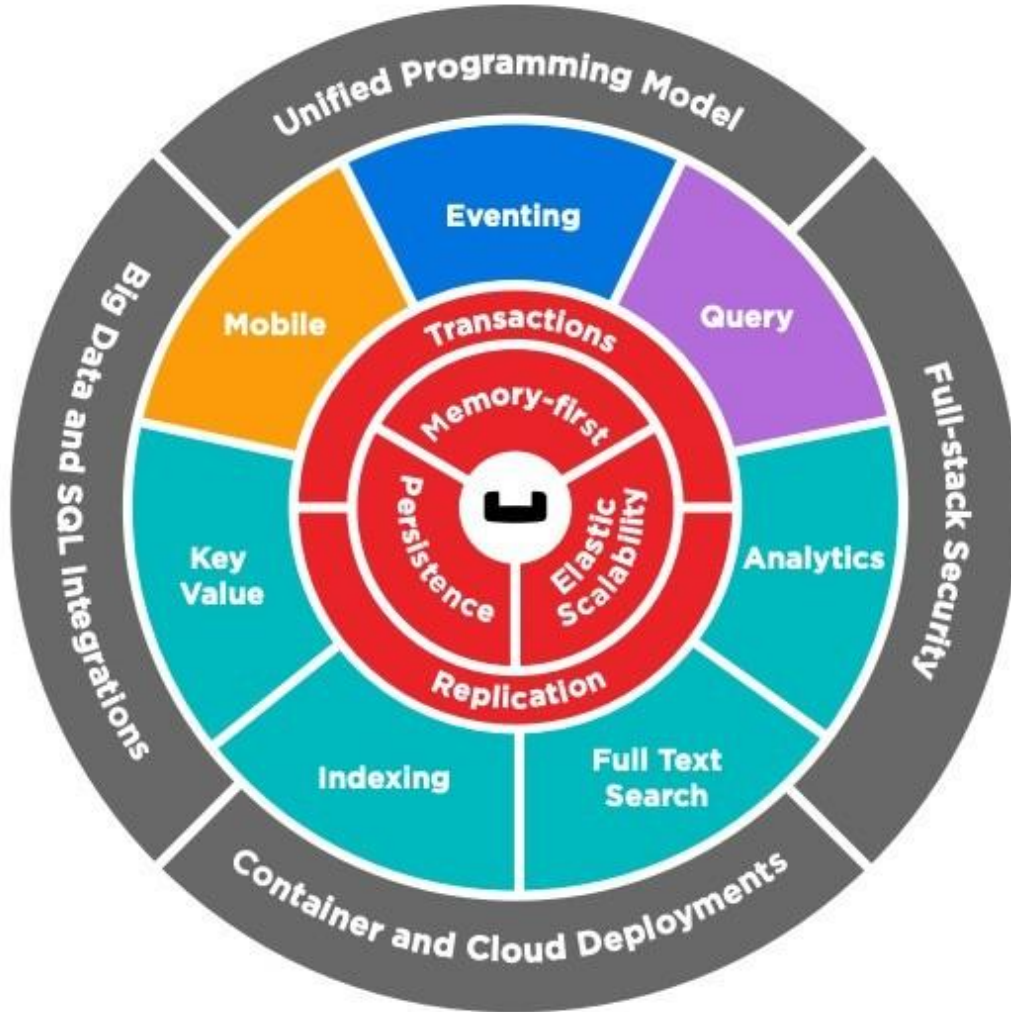


Data movement free from disk bottlenecks



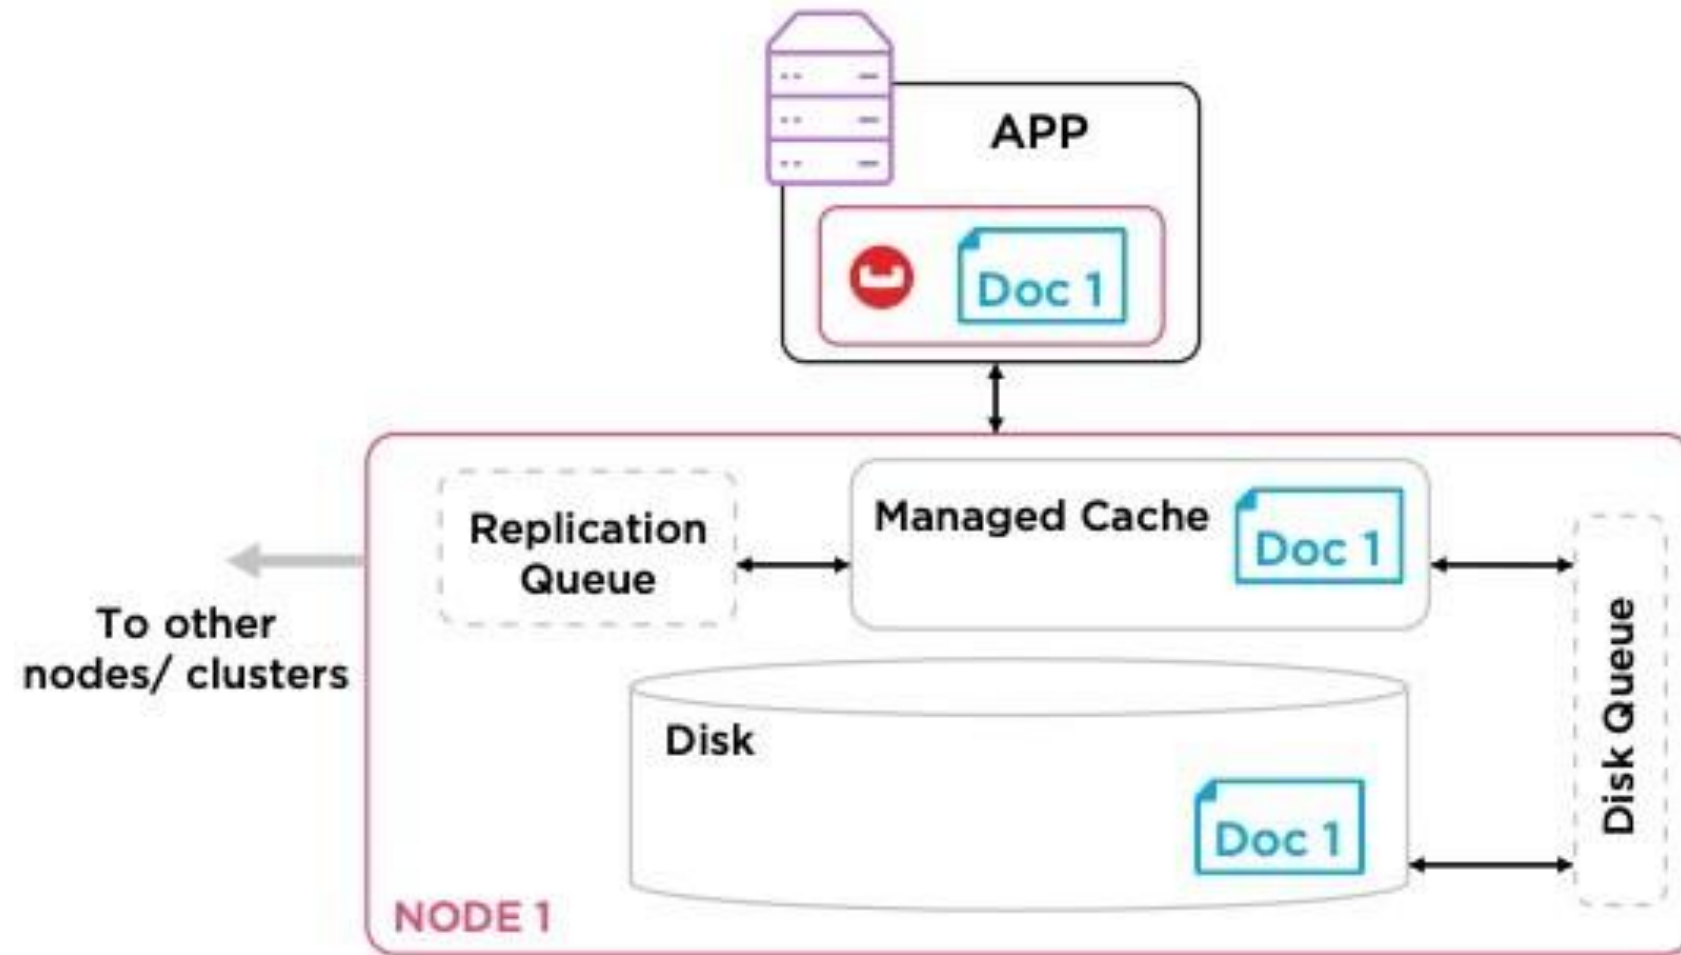
- In-memory streaming of updates to all components
- In-memory cache
- Memory-only data buckets
- Memory-only indexes

# Persistence



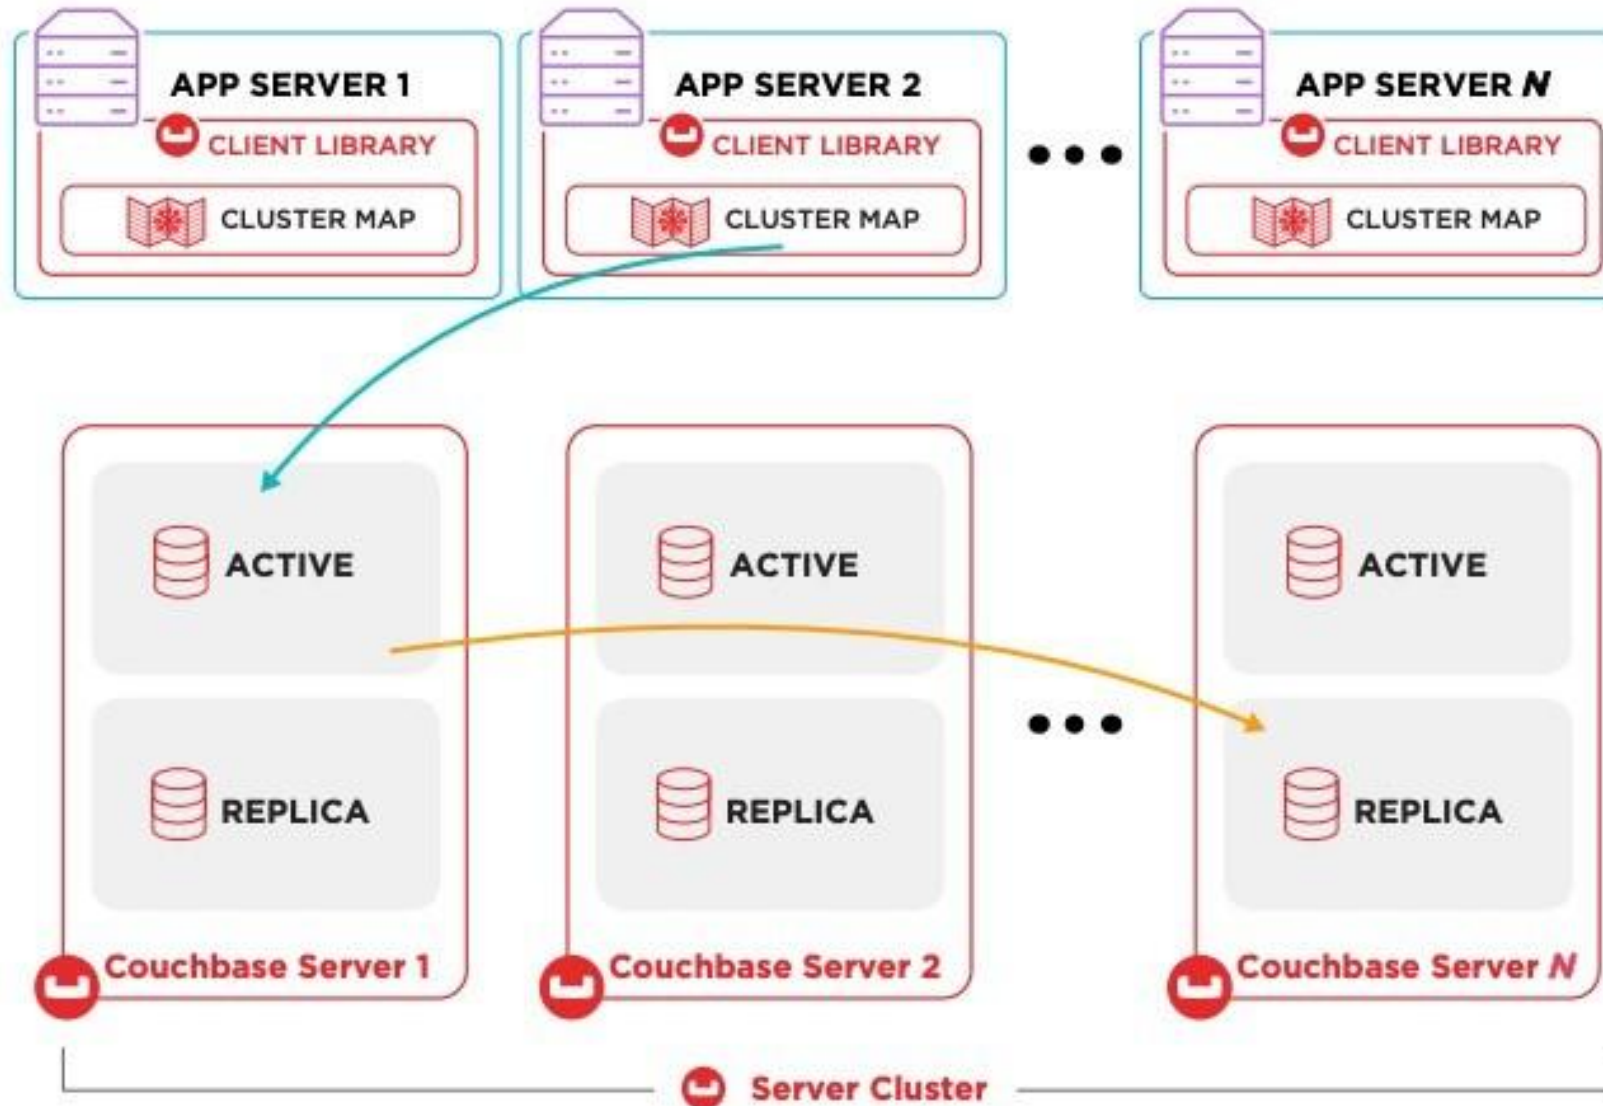


# Persistence




**Couchbase Server Asynchronous Architecture**

# Persistence



# Elastic Scalability



# Multi-Dimensional Scaling



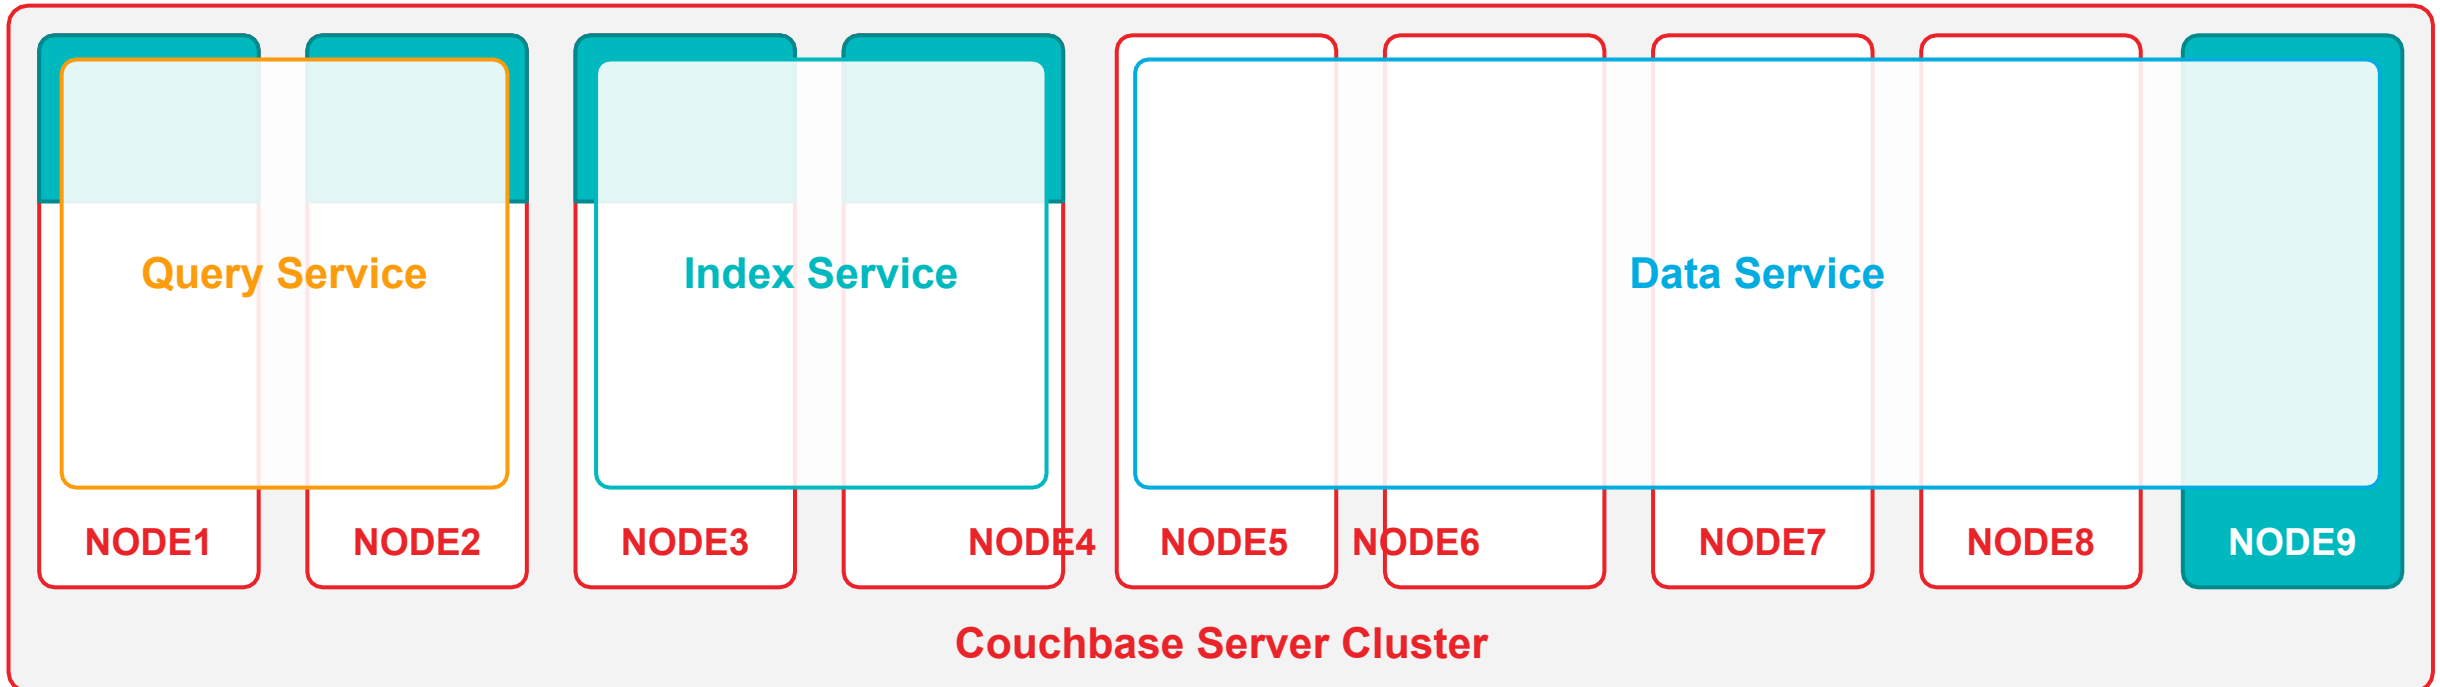
Independent Scalability for Best Computational Capacity – *per Service*

More CPU for Query Processing?

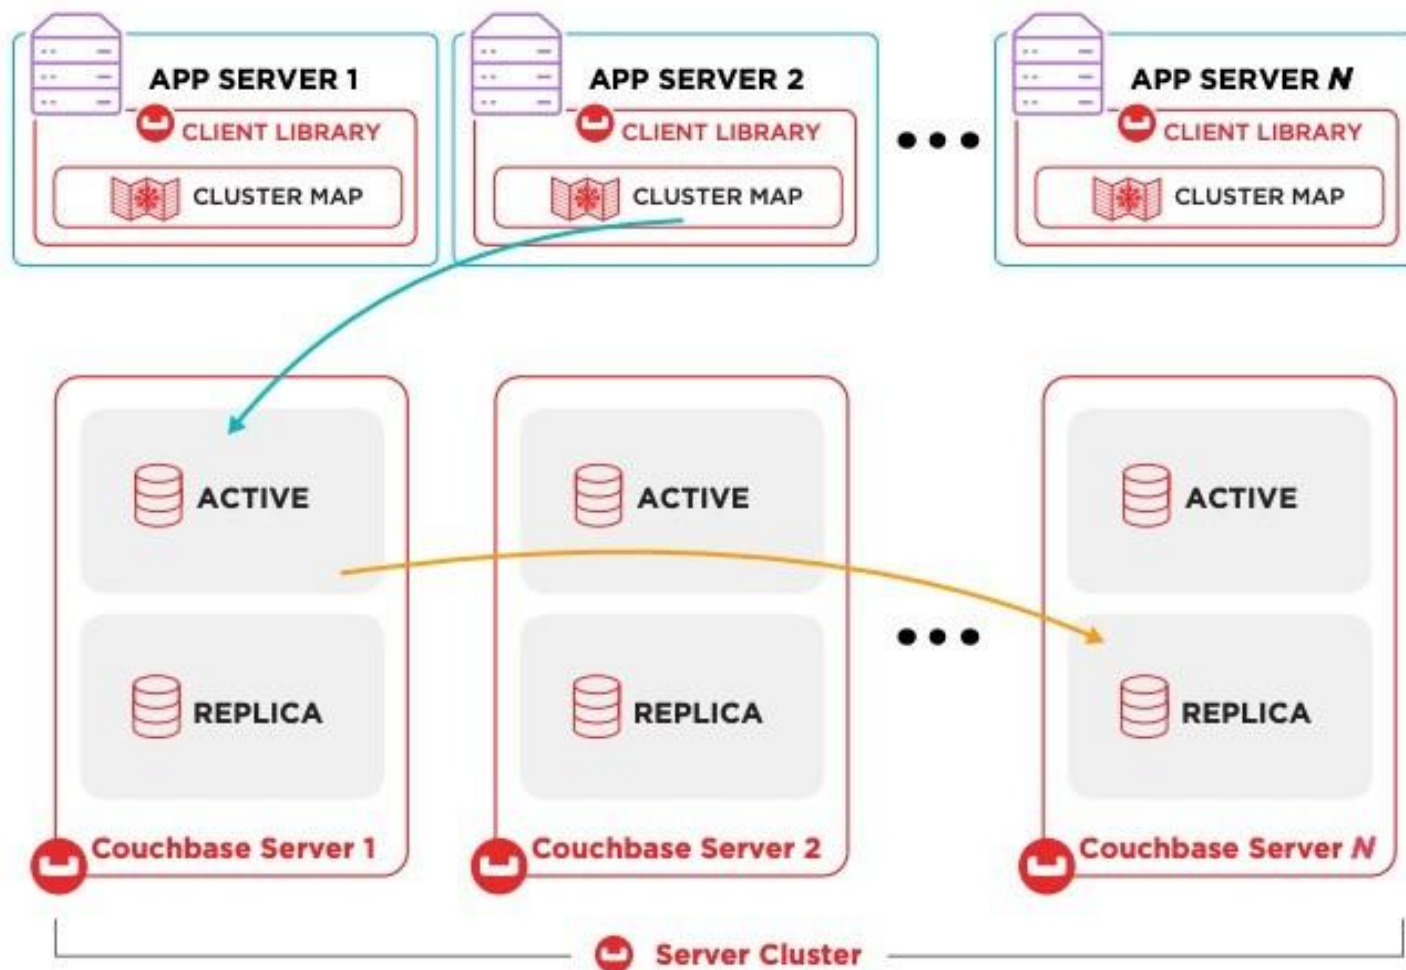
Heavier Indexing?

Scale up or out  
**Query Service** Nodes.

Scale up or out  
**Index Service** Nodes.




# Elastic Scalability & High Availability



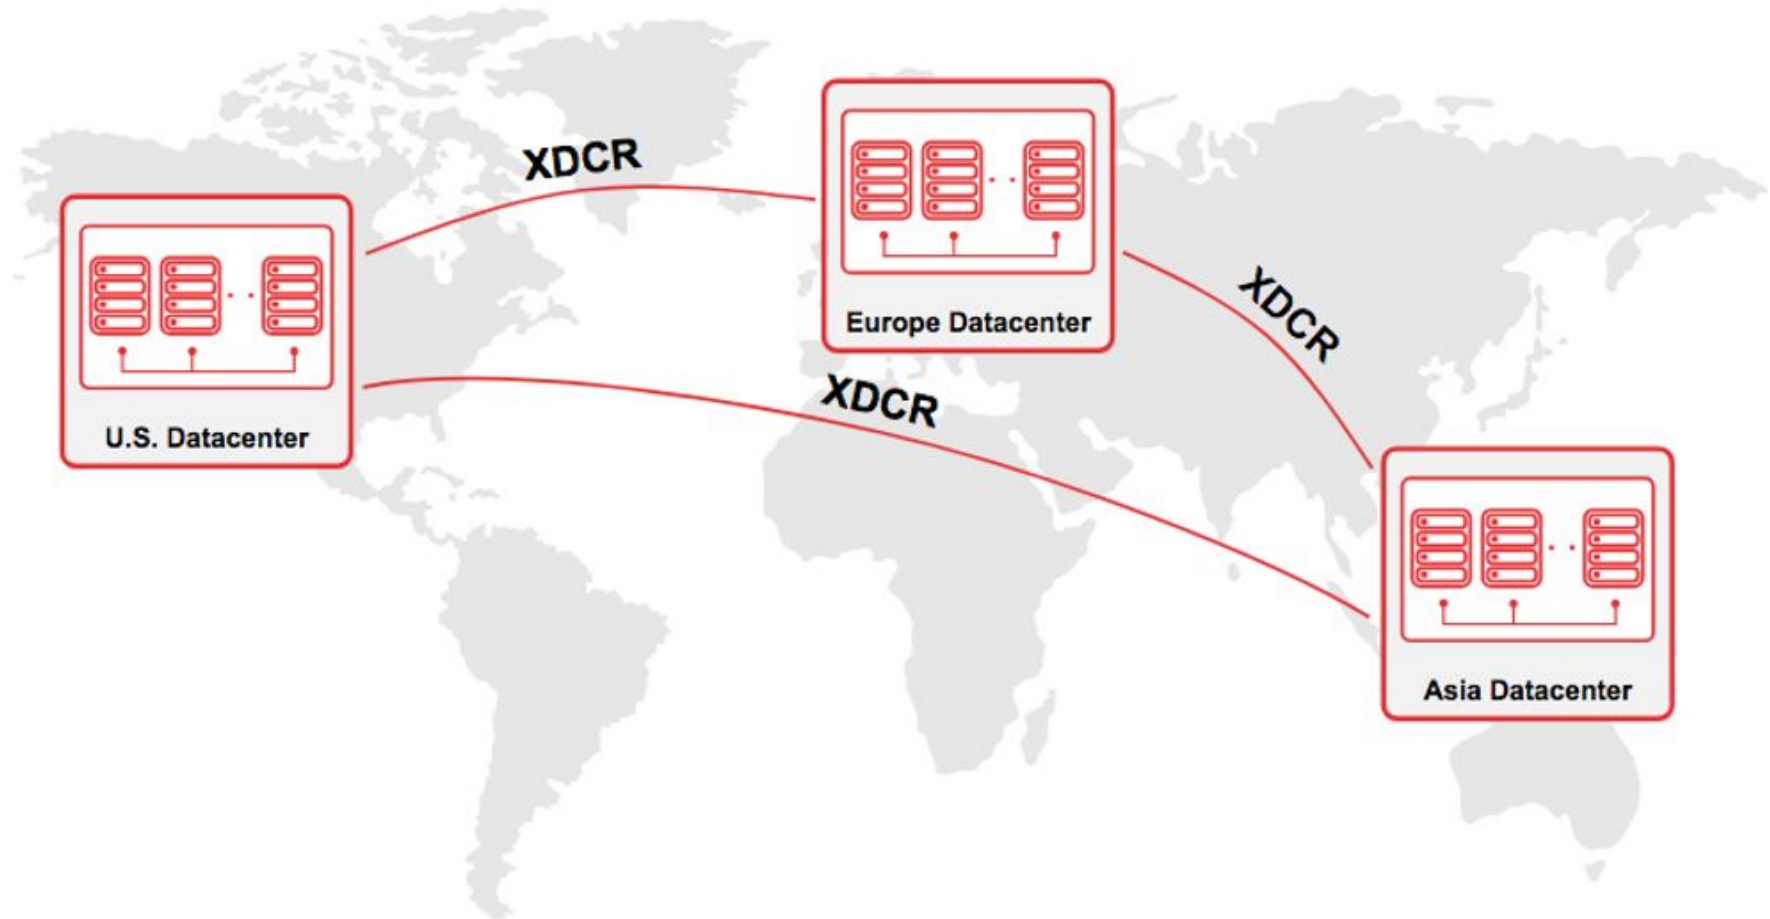
- **Scale on demand** with automatic partitioning and rebalancing
- **Build always-available apps** with memory-to-memory replication and automatic failover
- **Simplify development** with topology-aware clients and direct communication

# Replication






# Cross Data Center Replication (XDCR)



# Transactions





# Scopes and Collections = Schemas and Tables

---



Relational Model	Couchbase
Server	Cluster
Database	Bucket
<b>Schema</b>	<b>Scope</b>
<b>Table</b>	<b>Collection</b>
Row	Document (JSON or BLOB)




# Multi-Document ACID Transactions

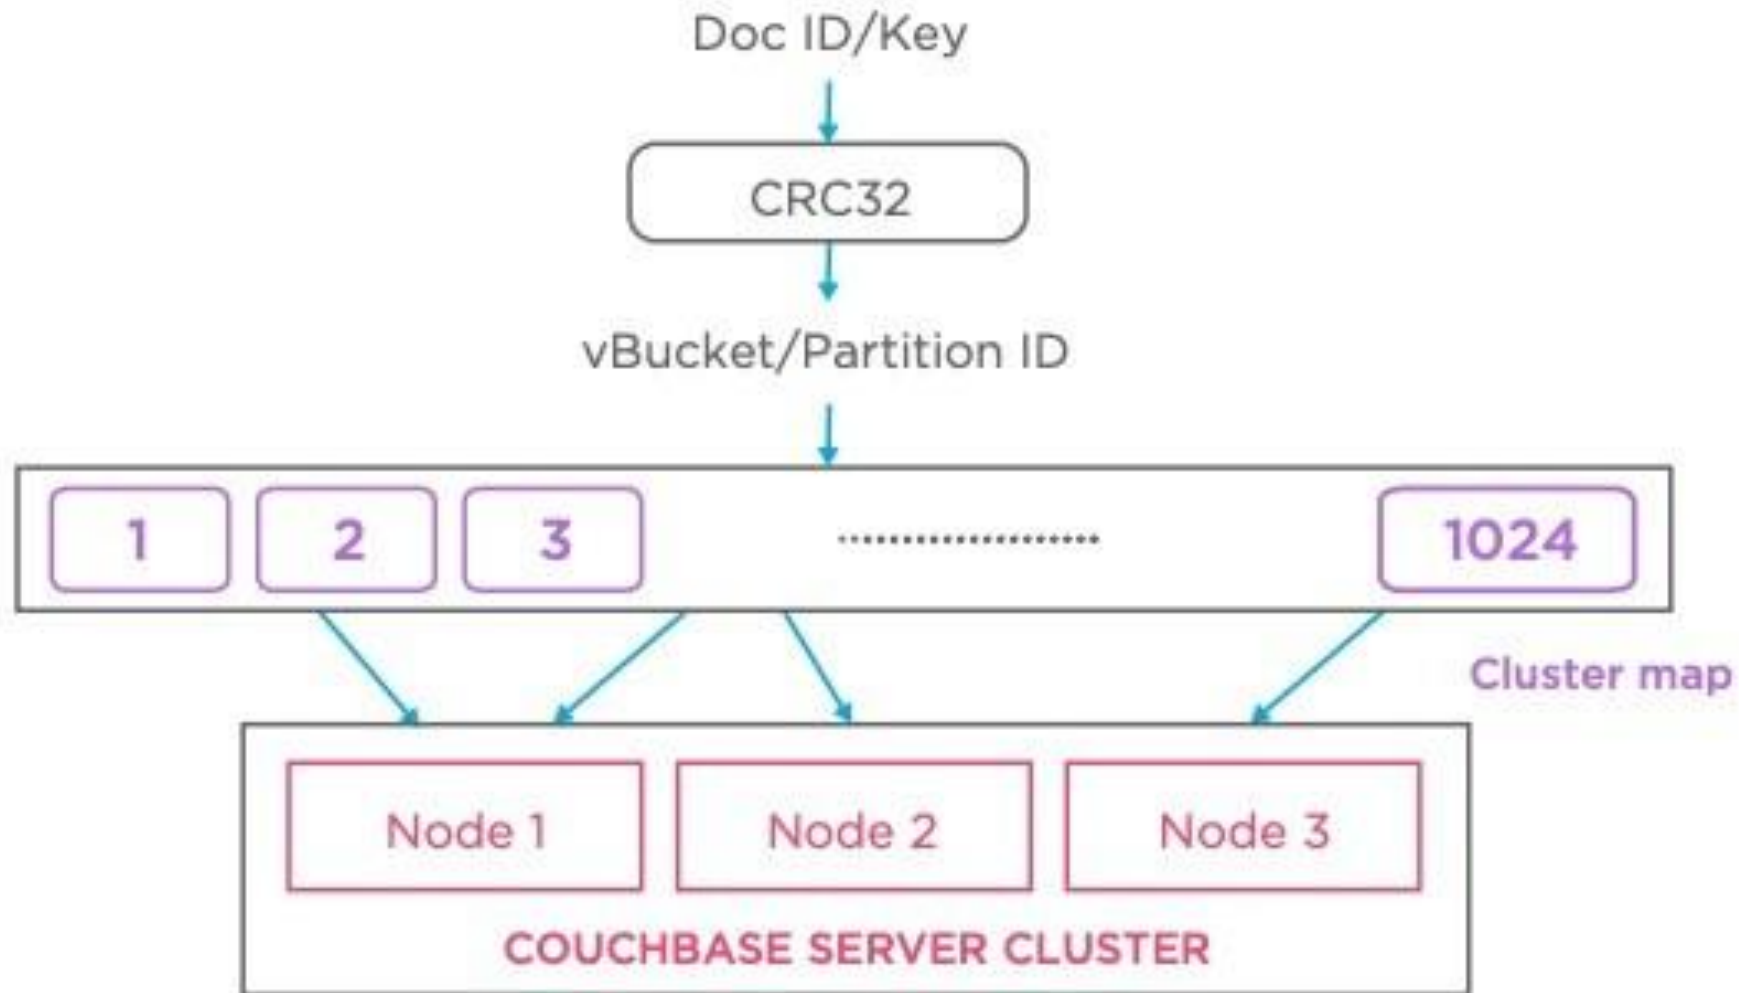
A	Atomicity	Guarantees all-or-nothing semantics for <b>updating multiple documents</b> in more than one shards on different nodes.
C	Consistency	<b>Replicas</b> immediately consistent with the transaction commit. <b>Indexes and XDCR</b> eventually consistent with the transaction commit (N1QL can enforce strong consistency upon read with request_plus)
I	Isolation	<b>Read Committed</b> isolation for concurrent transactions
D	Durability	<b>Data protection under failures</b> : 3 different levels - <i>replicate to majority of the nodes; replicate to majority and persist to disk on primary; or persist to disk on majority of the nodes.</i>

```
START TRANSACTION;  
UPDATE customer SET balance = balance + 100 WHERE cid = 4872;  
SELECT cid, name, balance from customer;  
UPDATE customer SET balance = balance - 100 WHERE cid = 1924;  
SELECT cid, name, balance FROM customer;  
COMMIT ;
```

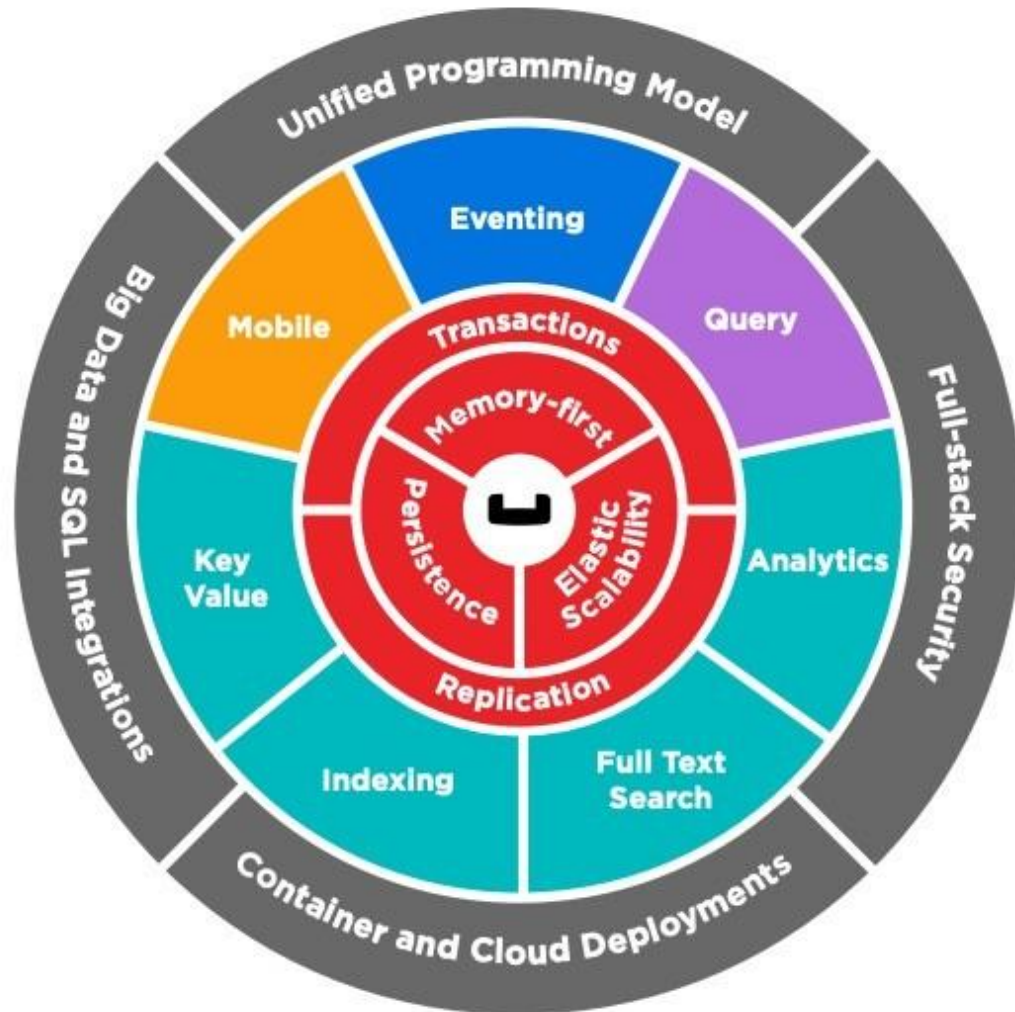
# Key / Value



# Key / Value Access



# Indexing / Query



# N1QL Origins



Named for 'Not 1<sup>st</sup> Normal Form Query Language'



First released with Couchbase 4.0, Oct. 2015



Work against flexible JSON Documents, not relational tables



Designed to be as close to SQL as possible



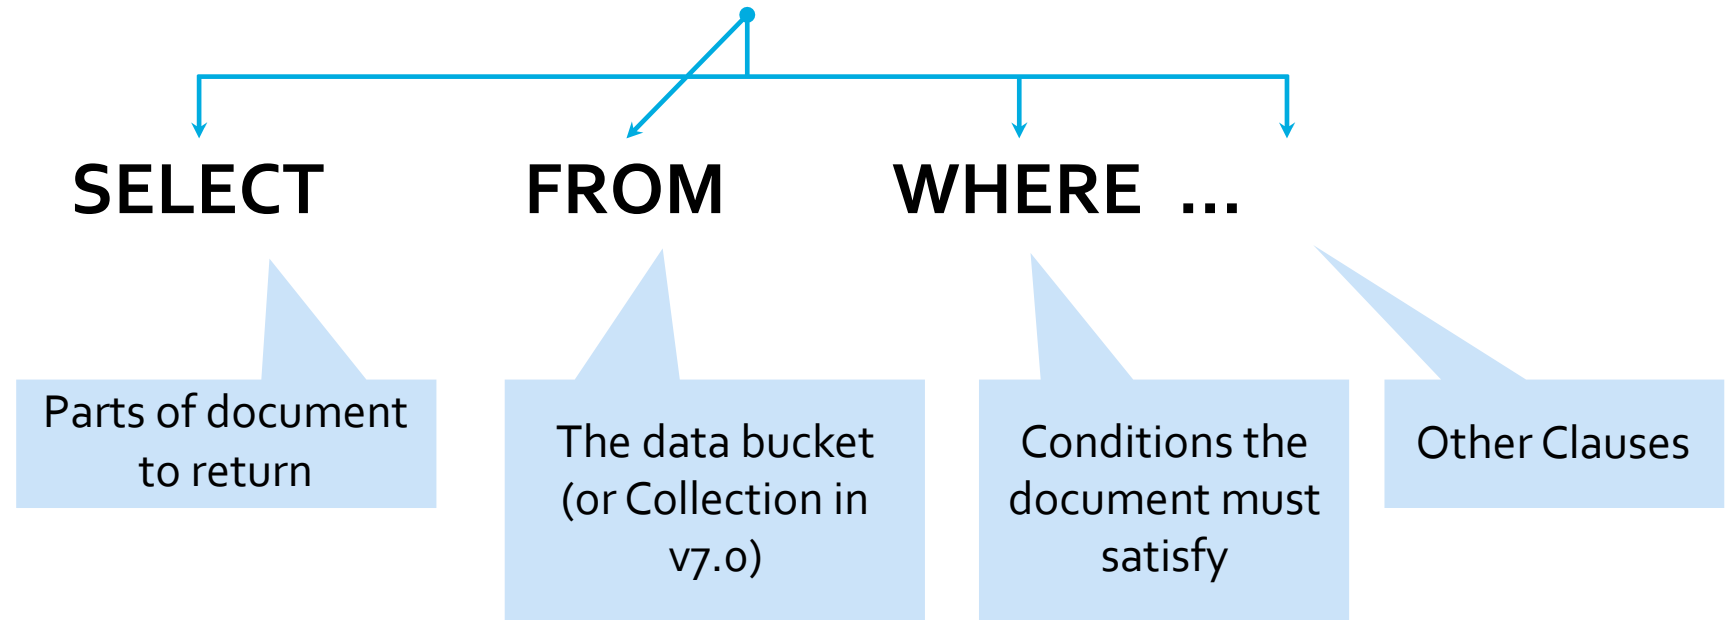
Overseen by Don Chamberlin, co-inventor of SQL



# Power of SQL

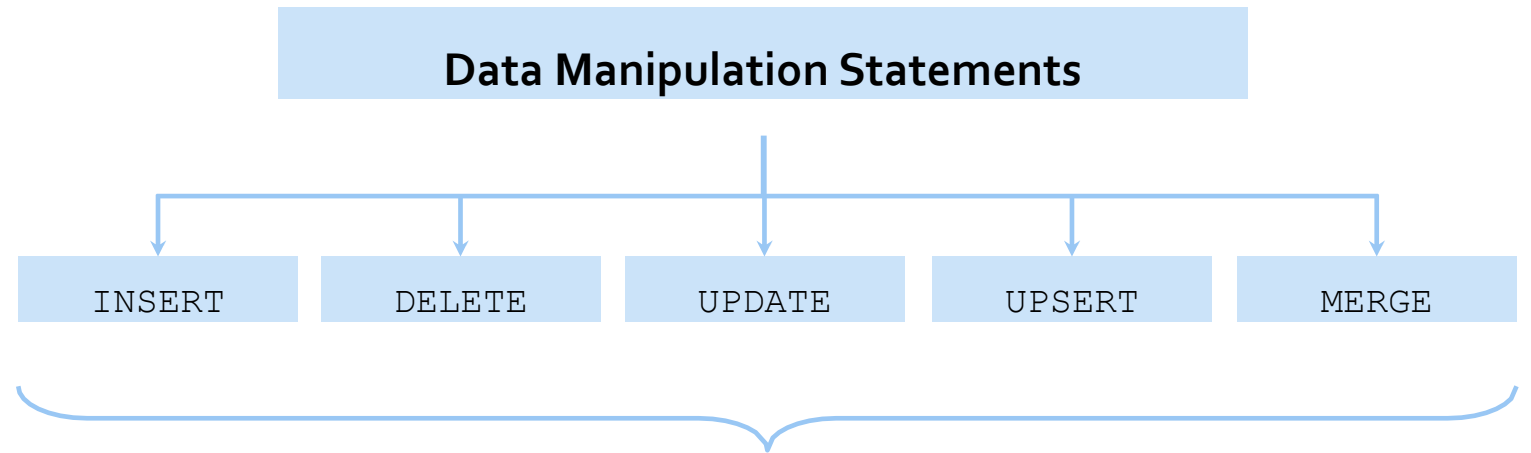
Looks familiar ?

## N1QL for Query



# Power of SQL

Includes DML



These statements allow you to create, delete, and modify the **data stored in JSON documents** by specifying and executing simple commands



# Power of SQL

## Key Features



Indexes

Use the CREATE INDEX and DROP INDEX statements to easily create, and delete indexes




Access Path

Traverse embedded documents with simple 'dot' notation



Aggregation

Provides aggregation operators, such as MIN, MAX, COUNT as well as grouping operators, the GROUP BY clause, and the group filter, HAVING



Joins

Retrieve data from multiple documents spanning one or more buckets



Nested Queries

Query within another query, embedded within the where clause

# Power of SQL

Extends SQL to work  
with Array

## Nesting

NEST lets you retrieve sub-records for a record



That is, for each left hand input, the matching right hand inputs are collected into an array, which is then embedded in the result

## Unnesting

UNNEST is the opposite of NEST



extracts nested sub-records from a record, making each extraction a separate record. Unnesting is performed on a single document

## Raw

SELECT RAW converts multiple results into an array (not a document)

# Power of SQL

## Support Joins

JOIN TYPE	Examples
LOOKUP JOIN	SELECT .. FROM customer c <b>JOIN</b> orders o <b>ON KEYS</b> ['..']
INDEX JOIN	SELECT .. FROM customer c <b>JOIN</b> orders o <b>ON KEYS</b> o.customer_id
ANSI JOIN	SELECT .. FROM customer c <b>JOIN</b> orders o <b>ON</b> o.customer_id = c.id
ANSI JOIN Complex	SELECT .. FROM `travel-sample` airline <b>JOIN</b> `travel-sample` route <b>ON</b> route.airlineid = "airline_"    toString(airline.id) <b>AND</b> route.type = "route"
ANSI JOIN with IN CLAUSE	SELECT .. FROM `travel-sample` route <b>JOIN</b> `travel-sample` airport <b>ON</b> airport.faa <b>IN</b> [ route.sourceairport, route.destinationairport ] <b>AND</b> airport.type = "airport"
ANSI LEFT OUTER	SELECT .. FROM `travel-sample` airport <b>LEFT JOIN</b> `travel-sample` route <b>ON</b> airport.faa = route.sourceairport <b>AND</b> route.type = "route"
ANSI JOIN with HASH JOIN	SELECT .. FROM `travel-sample` airport <b>JOIN</b> `travel-sample` route <b>USE HASH</b> (build) <b>ON</b> airport.faa = route.sourceairport <b>AND</b> route.type = "route"
ARRAY JOIN	SELECT .. FROM default b1 <b>JOIN</b> default b2 <b>ON</b> b2.c21 = b1.c11 <b>AND</b> b2.type = "right" <b>AND ANY</b> v <b>IN</b> b2.a21 <b>SATISFIES</b> v = b1.c12 <b>END</b>

# N1QL for Query - Example



## Movie

```
{
  "title": "The Avengers",
  "casting": [
    { "character": "Tony Stark / Iron Man",
      "name": "Robert Downey Jr",
      ...
    },
  ],
  "releaseYear": 2019,
  "revenue": 2792269749
}
```

## Person

```
{
  "birthYear": 1965,
  "deathYear": null,
  "primaryName": "Robert Downey Jr",
  "primaryProfession":
    "actor, producer, soundtrack"
}
```

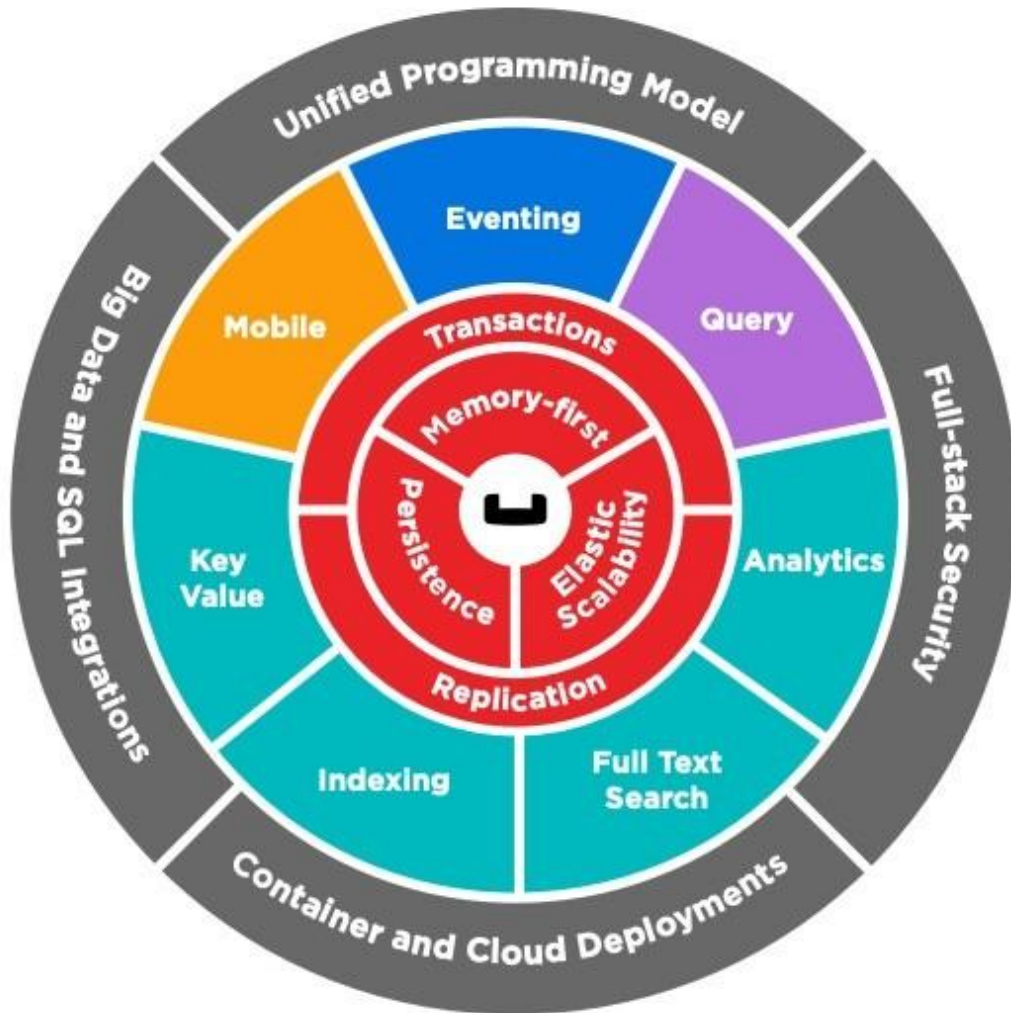
*Identify for each movie where Robert Downey Jr. has been playing in who is the youngest actor in the movie, organised by year.*

```
SELECT
  MIN(person.birthYear) AS year,
  movie.title
FROM moviedata AS movie
UNNEST movie.casting AS actor
JOIN moviedata AS person
  ON actor.name = person.primaryName
  AND person.type = "person"
WHERE movie.type = "movie"
AND ANY member IN movie.casting SATISFIES
  member.name = "Robert Downy Jr" END
GROUP BY movie.title
ORDER BY year DESC
```

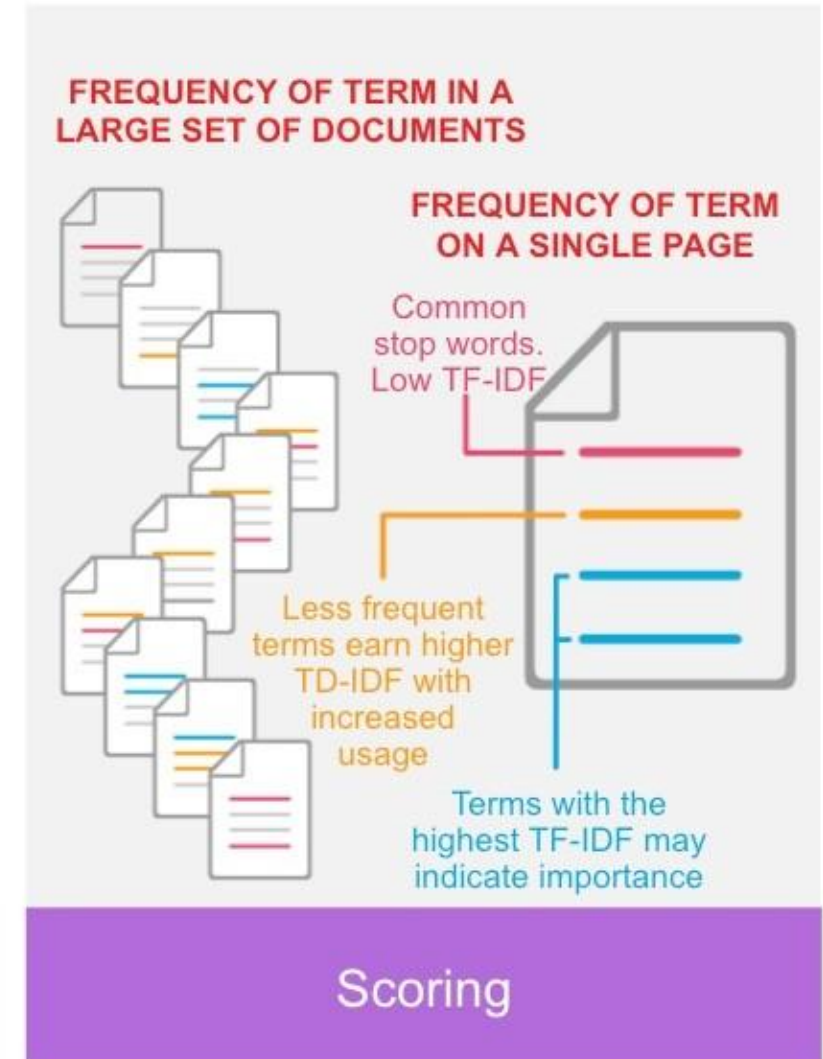
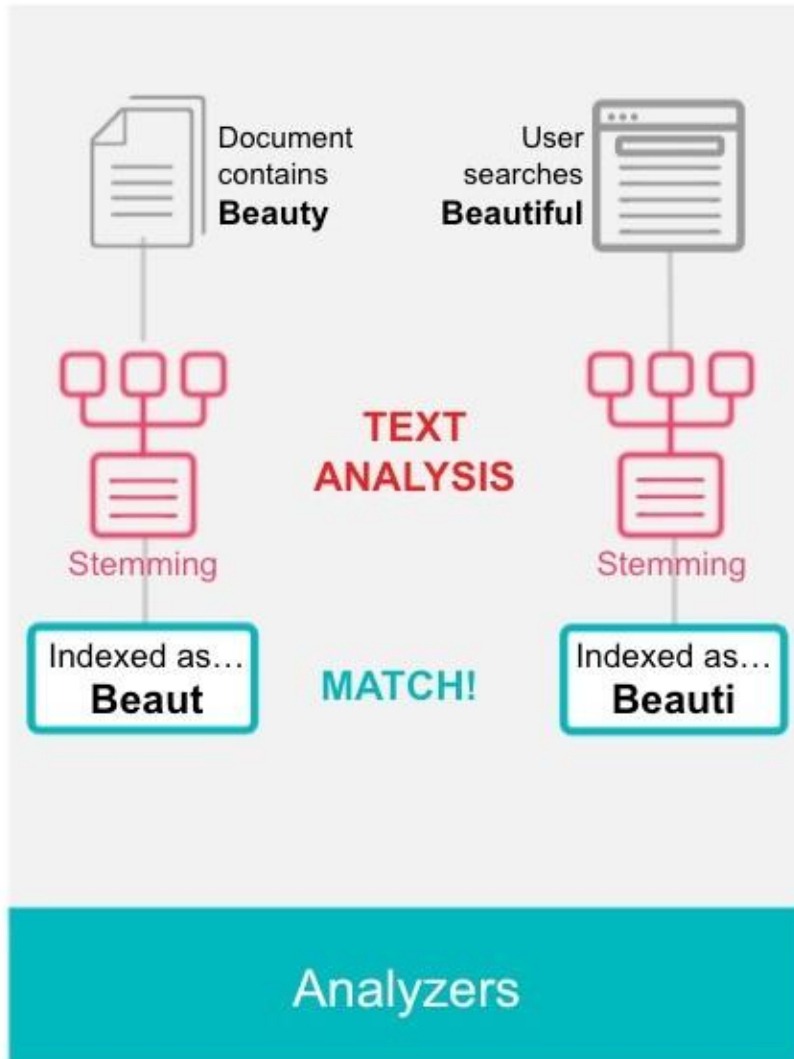
## Result

```
[
  {
    "title": "Charlie Bartlett",
    "year": 1978
  },
  {
    "title": "The Avengers",
    "year": 1973
  },
  {
    "title": "Captain America: Civil War",
    "year": 1973
  },
  {
    "title": "A Guide To Recognizing
      Your Saints",
    "year": 1973
  },
  ...
]
```

# Full Text Search



# Full Text Search




# Full Text Search



Simple Queries



Compound Queries



Range Queries (string, date, numeric)



String Queries (natural language)



Geospatial Queries



Non-analytic (i.e. exact match)




Special queries (for dev purpose)

Name	Language
ar	Arabic
bg	Bulgarian
ca	Catalan
cjk	Chinese   Japanese   Korean
ckb	Kurdish
da	Danish
de	German
el	Greek
en	English
es	Spanish (Castilian)
eu	Basque
fa	Persian
fi	Finnish
fr	French
ga	Gaelic
gl	Spanish (Galician)
hi	Hindi
hu	Hungarian

Name	Language
eu	Basque
fa	Persian
fi	Finnish
fr	French
ga	Gaelic
gl	Spanish (Galician)
hi	Hindi
hu	Hungarian
hy	Armenian
id, in	Indonesian
it	Italian
nl	Dutch
no	Norwegian
pt	Portuguese
ro	Romanian
ru	Russian
sv	Swedish
tr	Turkish

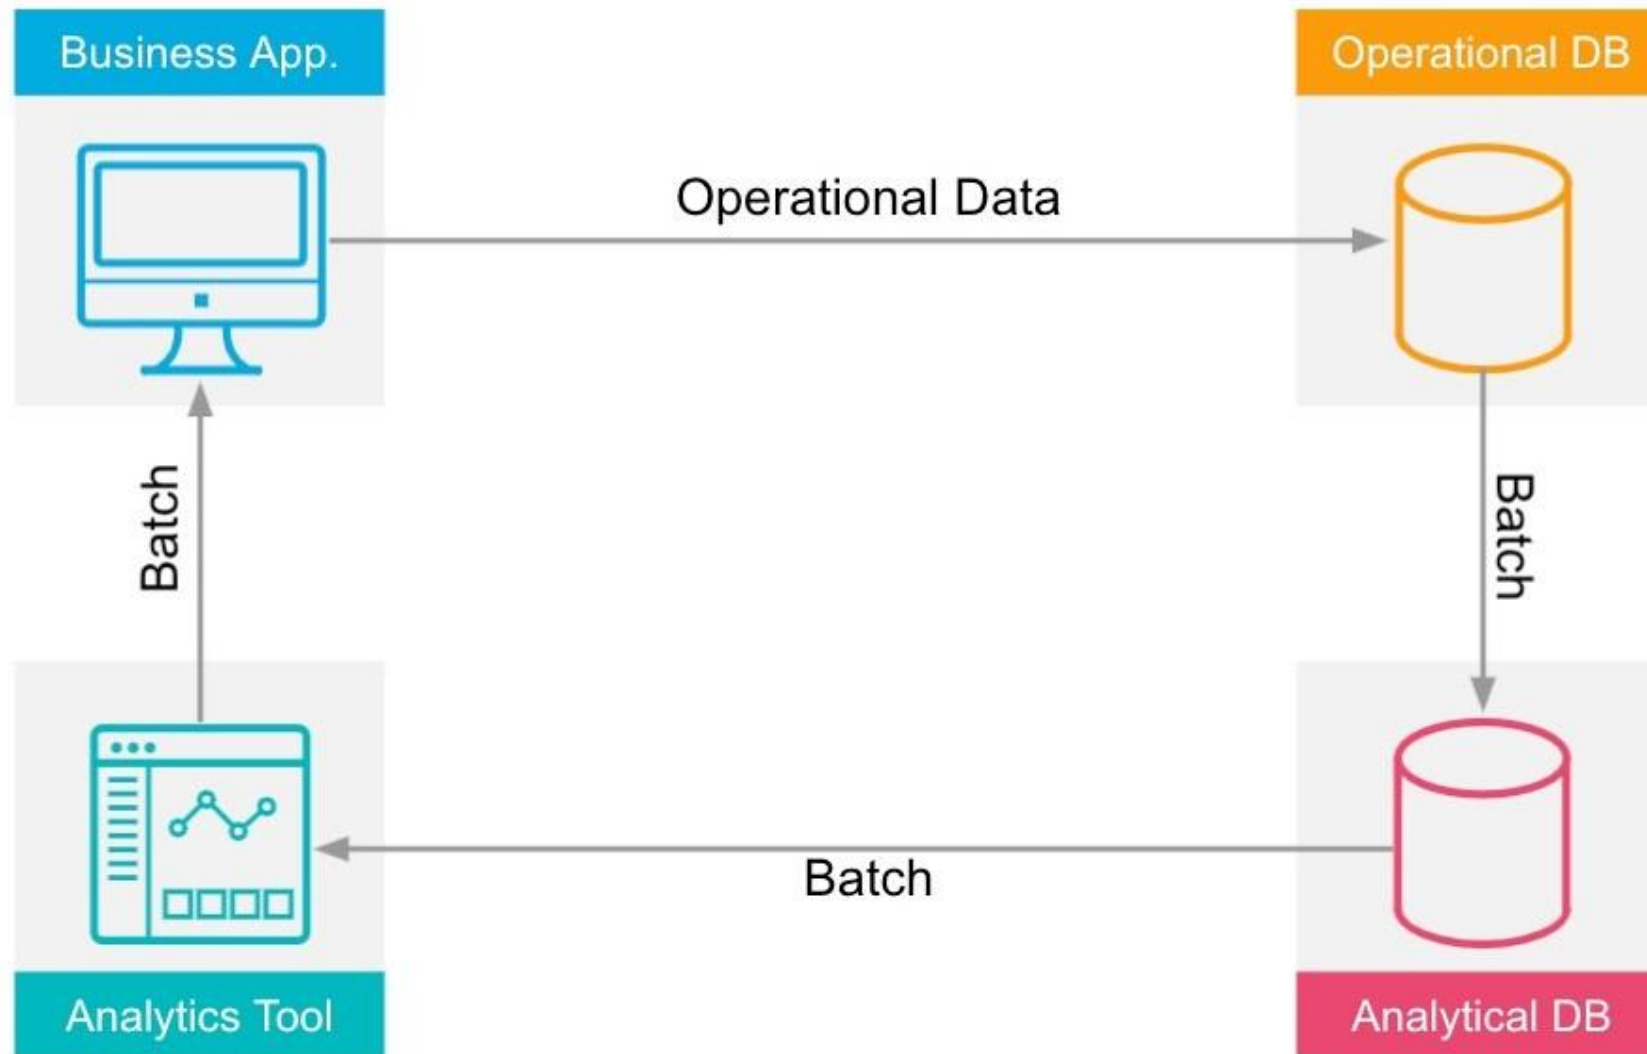


# Analytics

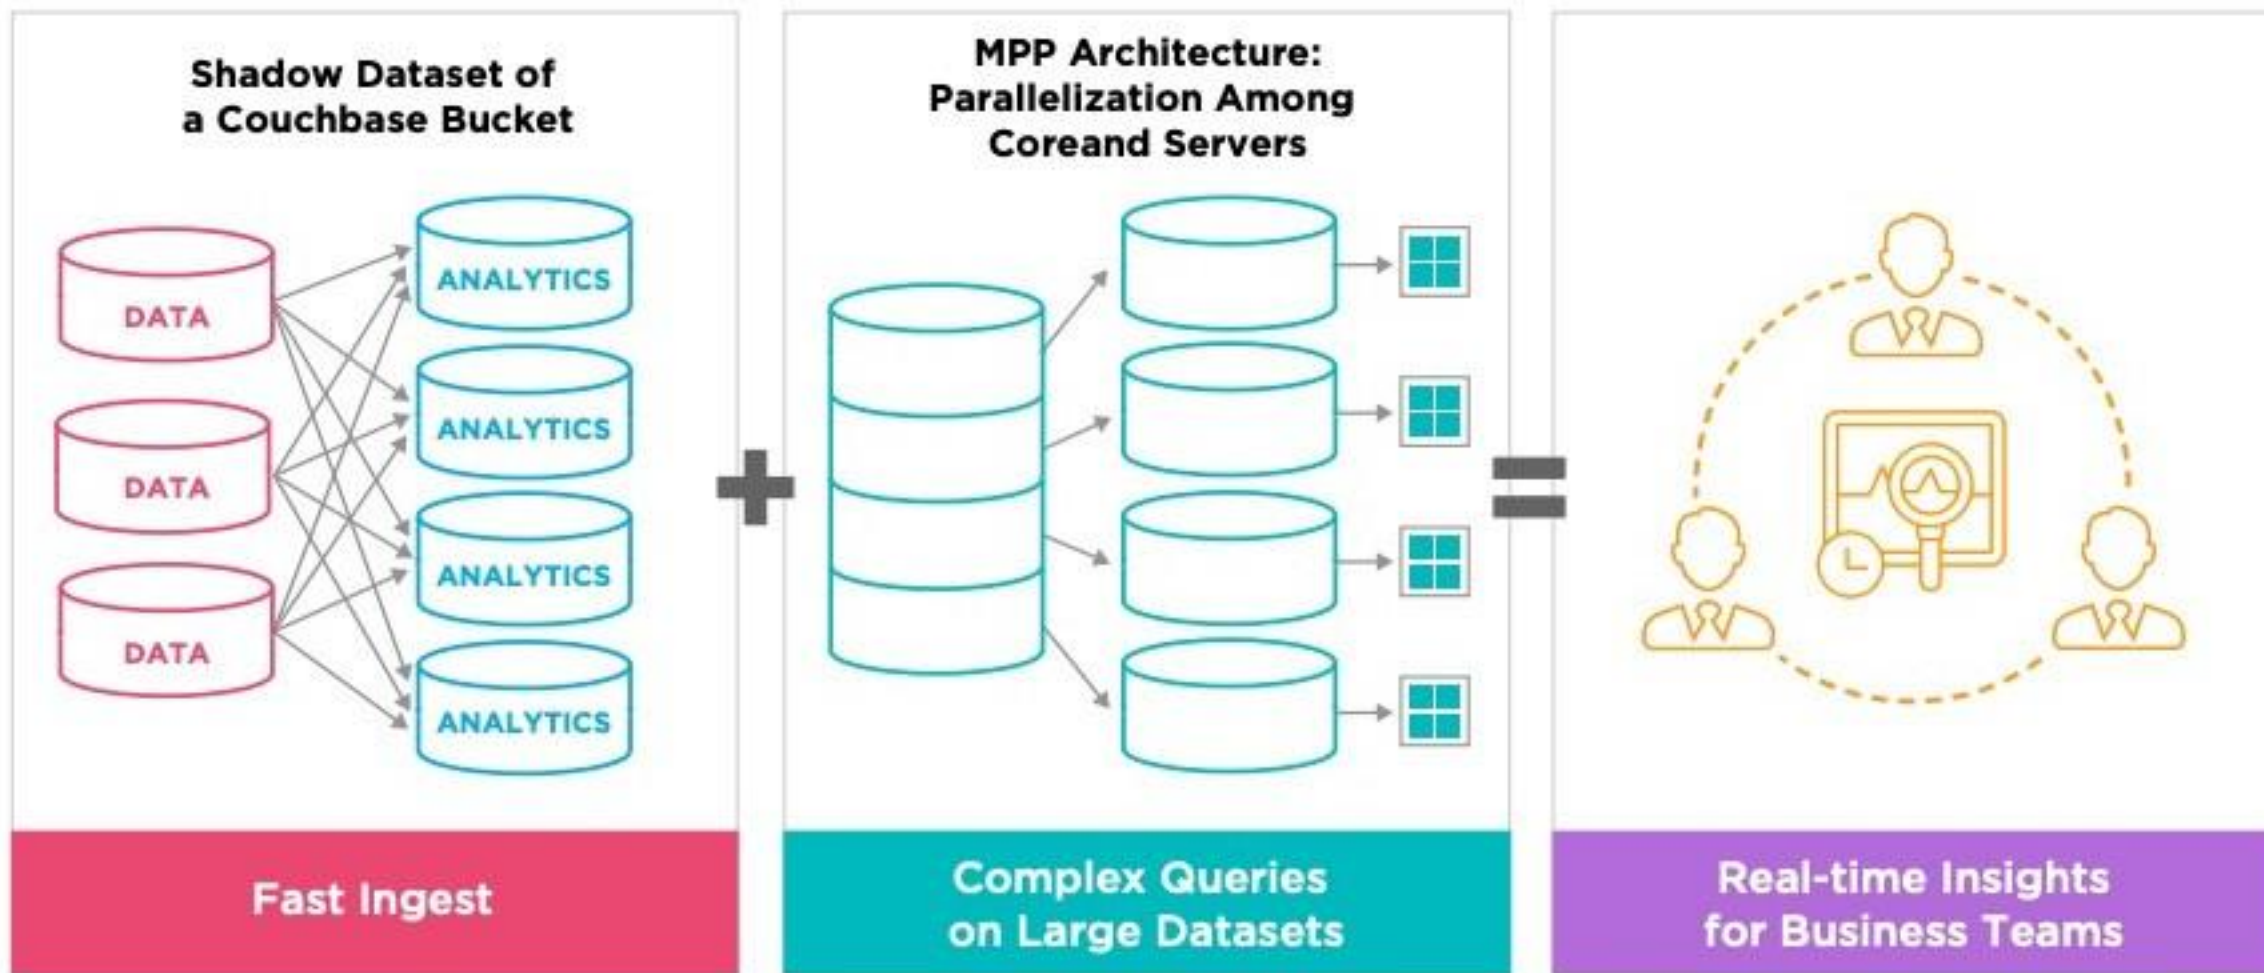





# Analytics: A typical setup with ETL



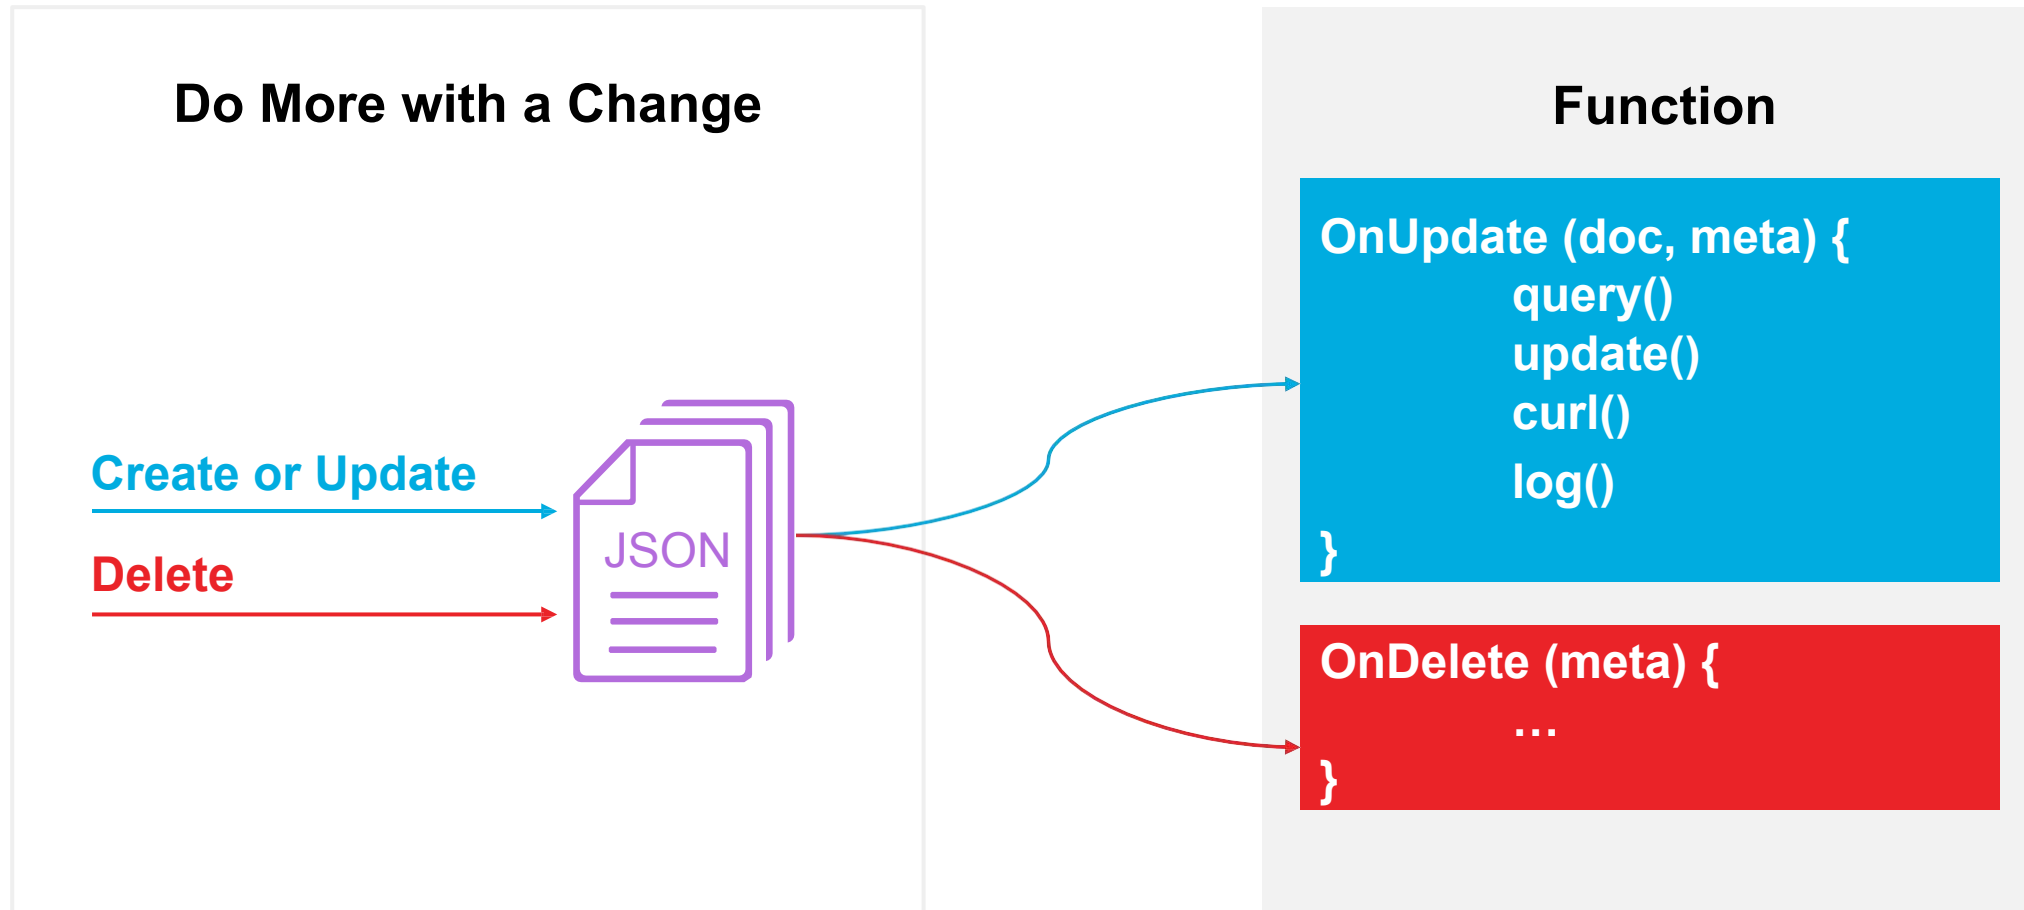
# Analytics on Couchbase: NoETL



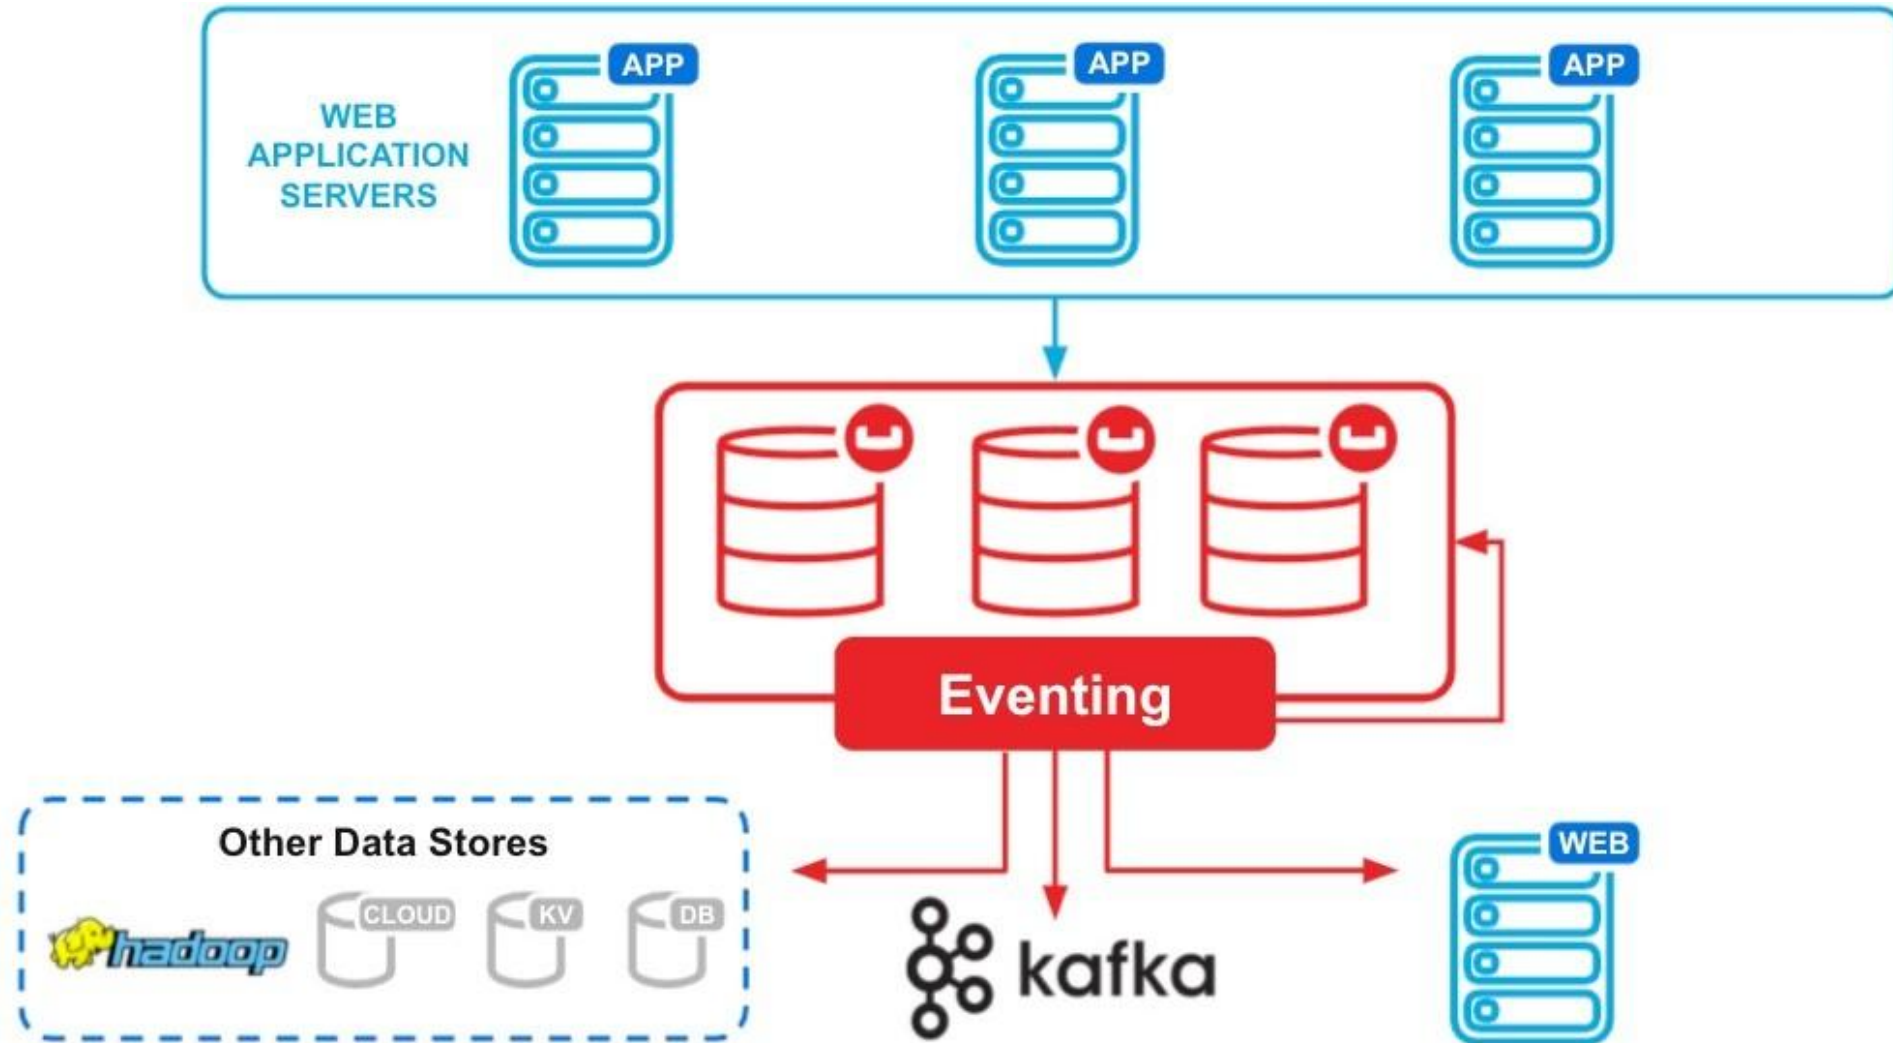
# Eventing




# Eventing



# Eventing

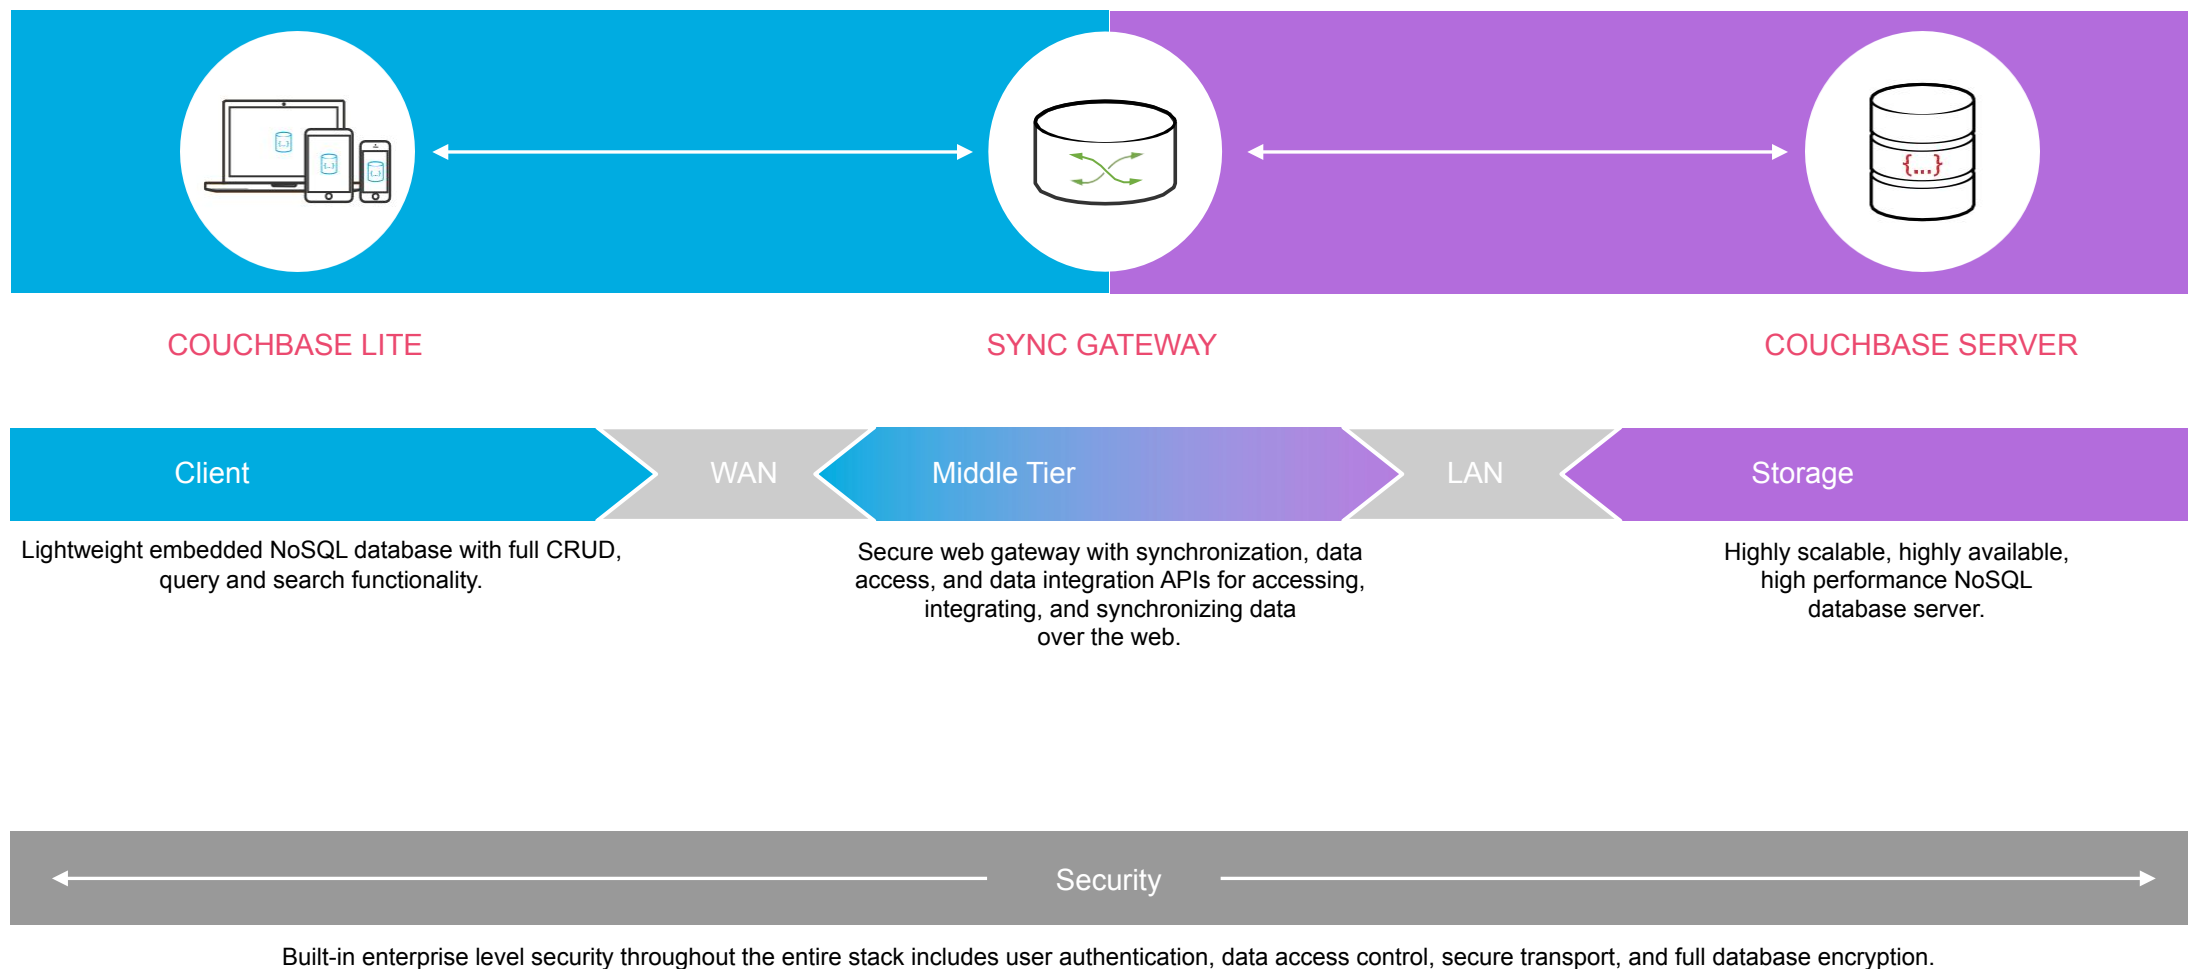


# Mobile

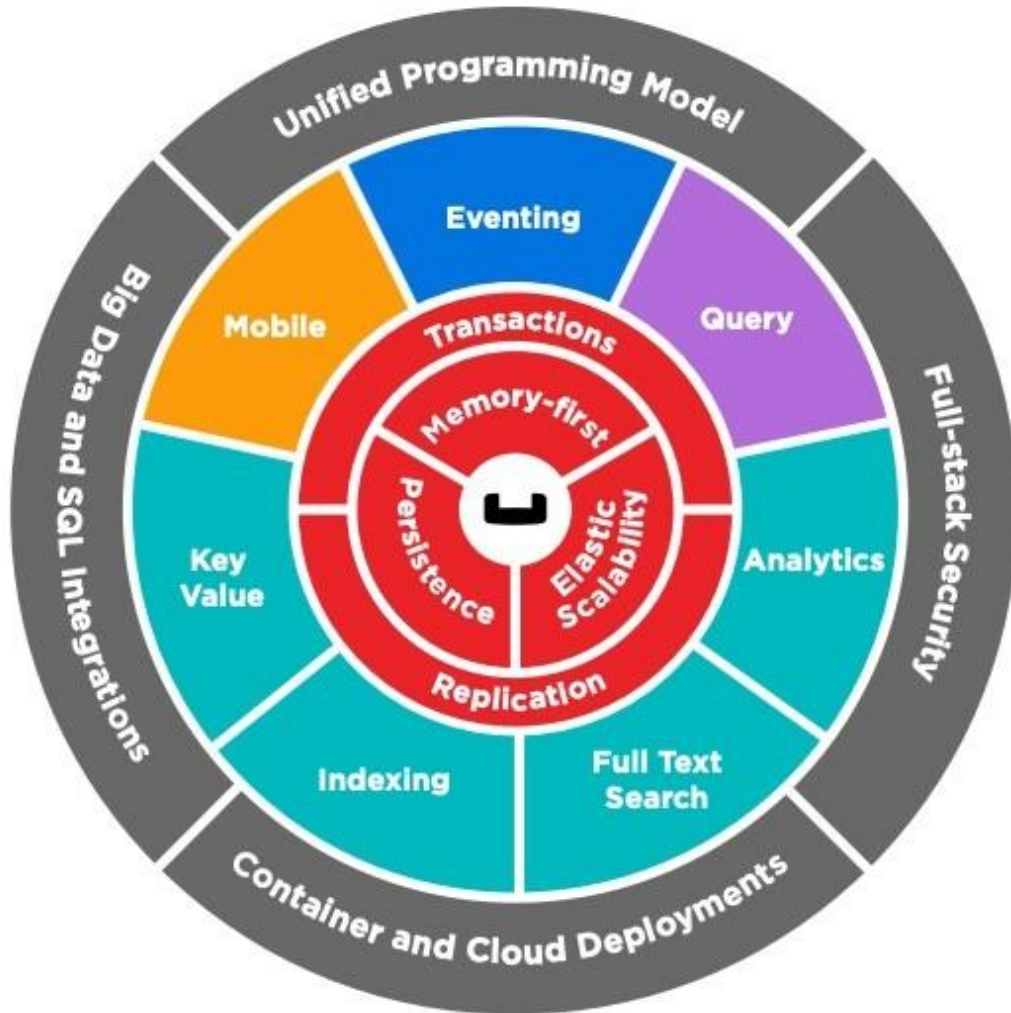




# Mobile



# Integrations






# Big Data and SQL Integrations


---



# Unified Programming Model



# Unified Programming Model – Single API



# SDKs

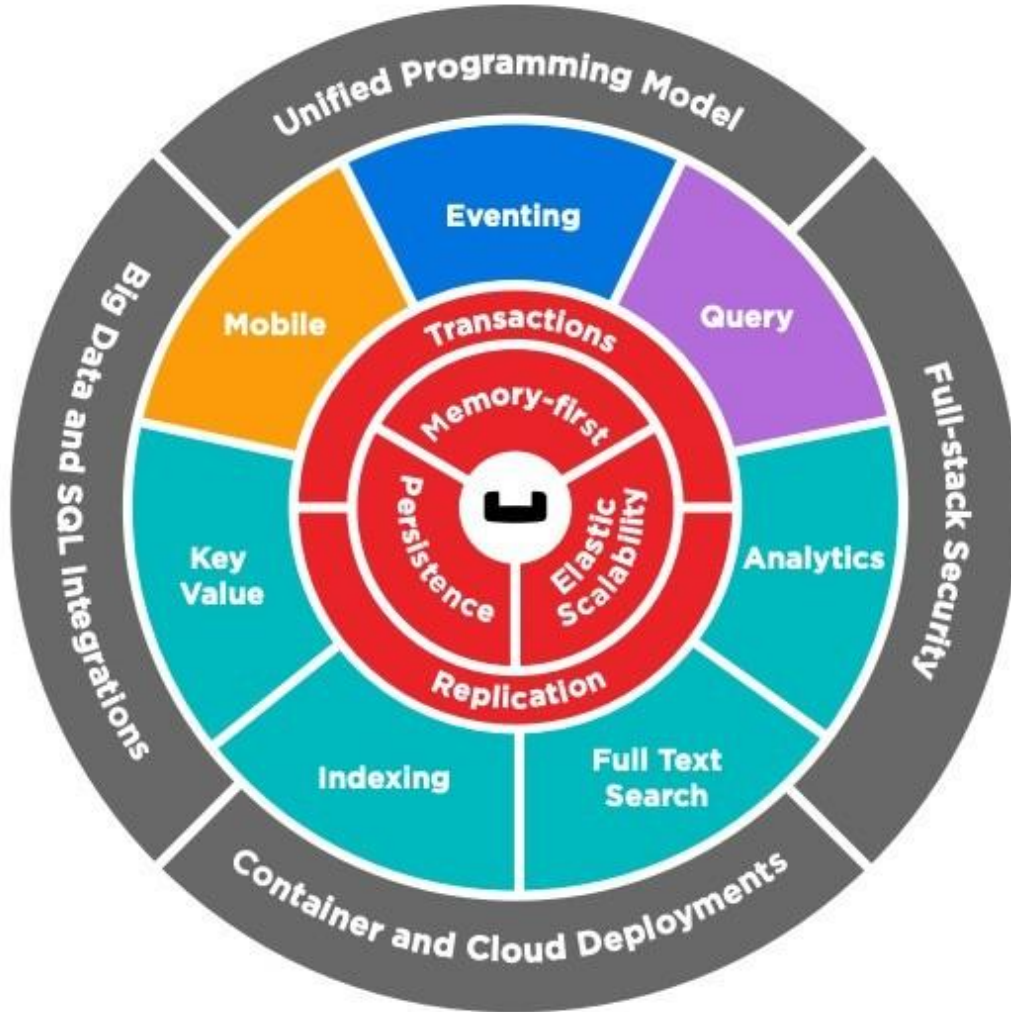
---



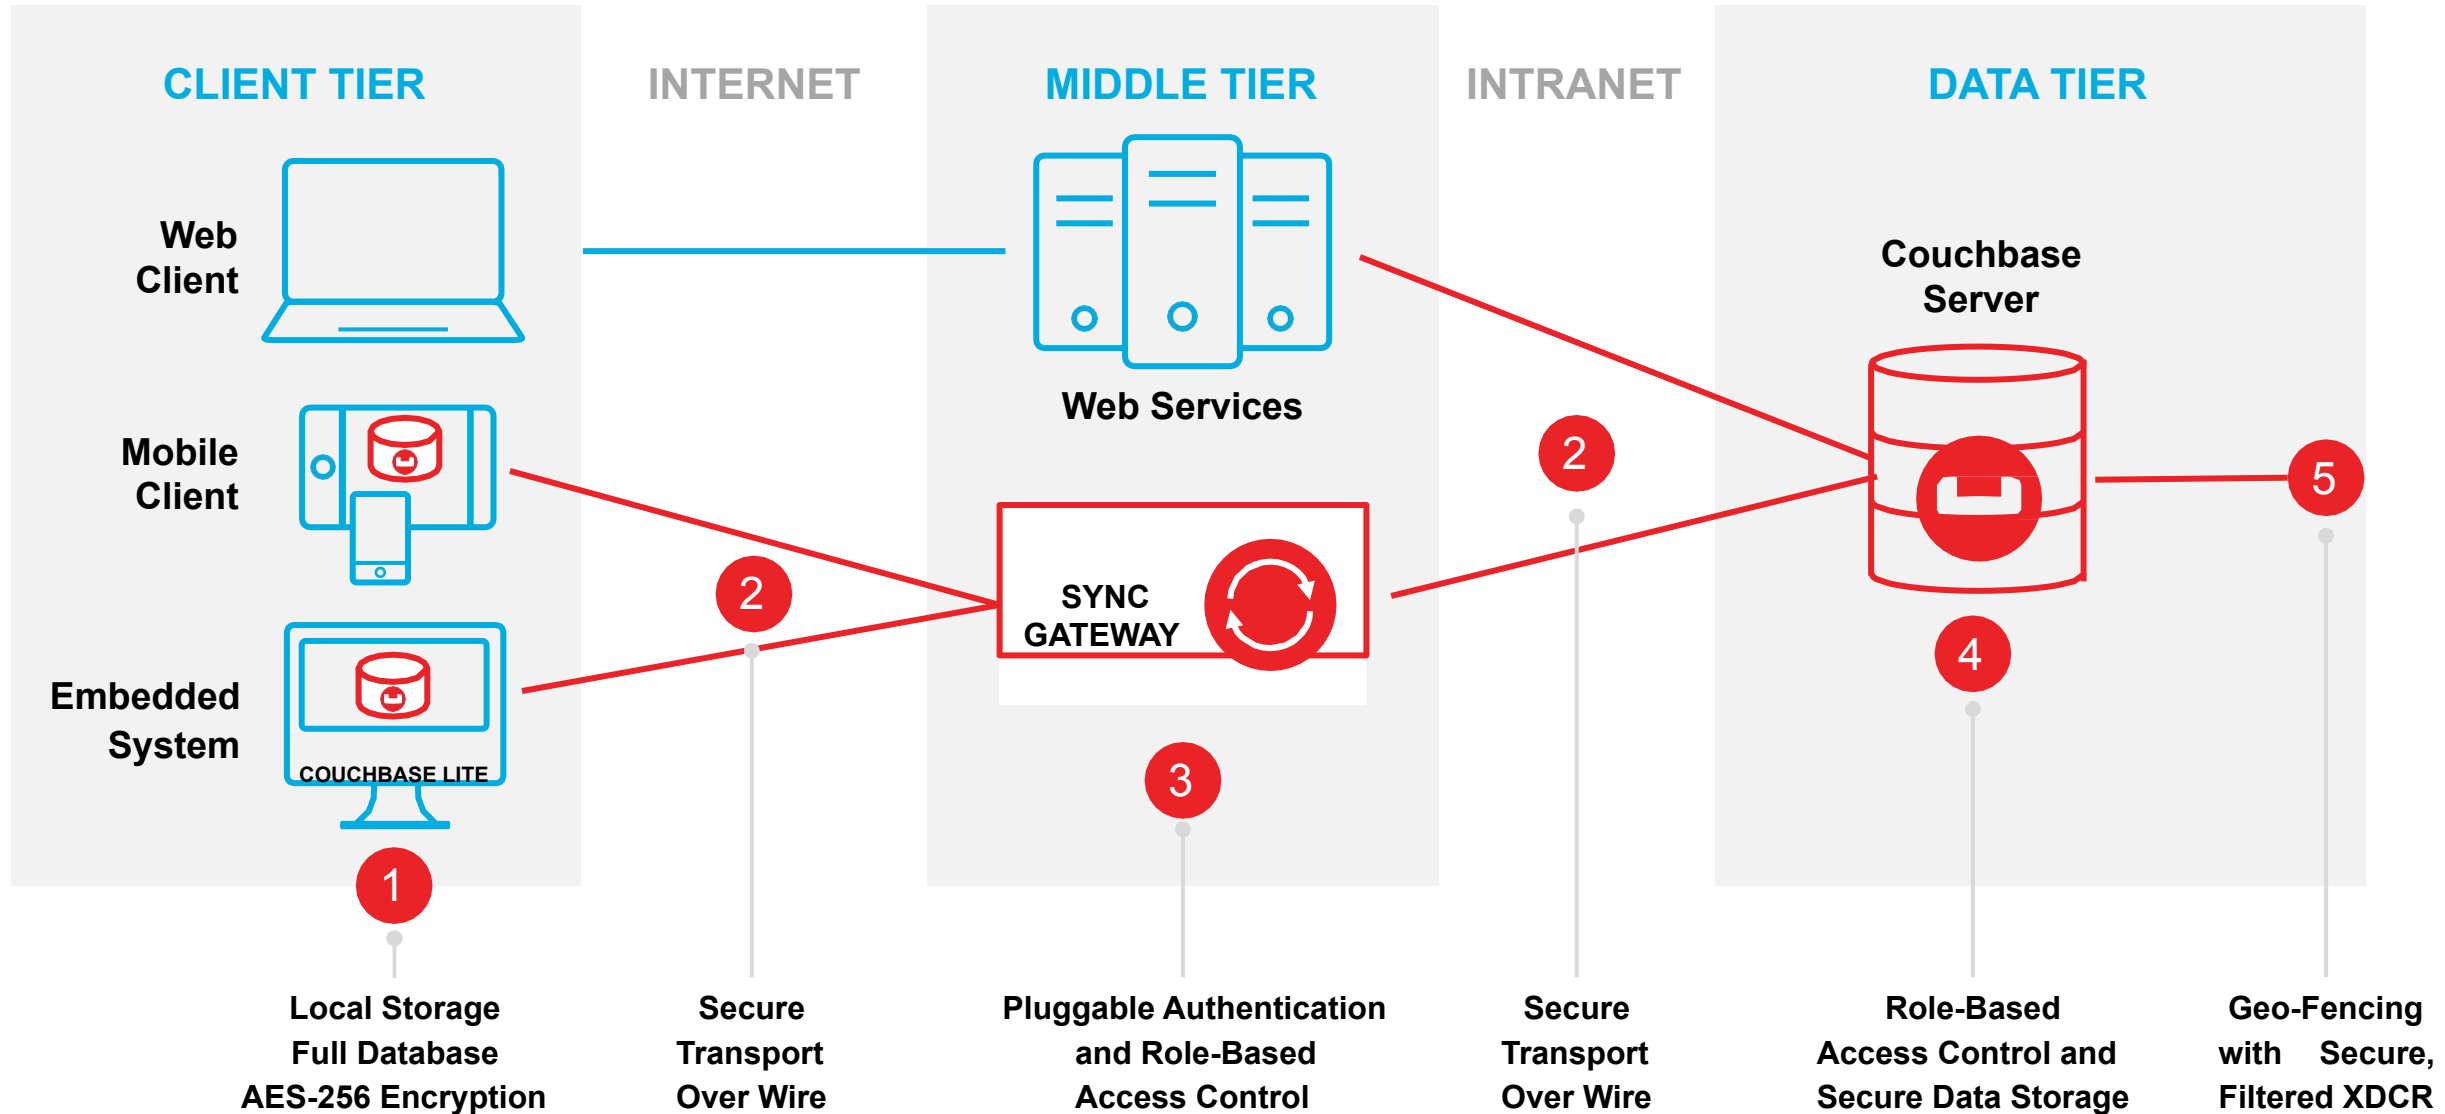
Community Support



# Security



# Full Stack Security






# Full Stack Security



Authentication	Authorization	Crypto	Auditing	Operations
App/Data: SASL AuthN Admin: Local or LDAP Users or LDAP Groups PAM Authentication	Local Admin User Local Read-Only Admin RBAC for Admins RBAC for Applications (since 5.0)	TLS admin access TLS client-server access Secure XDCR X.509 certificates for TLS Data-at-rest Encryption* Field-level Encryption (since 5.5) Secret Management Support for Configurable TLS Cipher Suites	Admin auditing API request auditing (since 5.5) N1QL auditing (since 5.5)	Security management via UI/CLI/REST

# Deployments





# Deploy Anywhere



## Database-as-a-Service



- Fully-Managed service
- Automated and Self-Service
- Agile and flexible – deploy in minutes
- Highly available, reliable and secure
- Comprehensive monitoring and support
- Eliminate operational overhead

**Couchbase Cloud** Public Clouds

## Self-Managed on IaaS



- Customer managed and operated
- Deployed in the Public Cloud
- Requires DBA's and Infrastructure admins from customer to manage
- Kubernetes Autonomous Operator and deployment templates speed deployment and reduce management burden

**Public Clouds**

## Self-Managed On-Premise



- Customer managed and operated
- Maximum flexibility, customizability and performance
- Requires DBA's and Infrastructure admins to manage
- Kubernetes Autonomous Operator speeds deployment and reduces management burden

**Private Clouds**


**FLEXIBLE DEPLOYMENT OPTIONS TO MEET YOUR REQUIREMENTS**

# COUCHBASE ARCHITECTURE

ANY CLOUD,  
ANY DEVICE,  
MANY  
SERVICES

## ONE data platform

Combines multiple **natively developed services** **Infrastructure-agnostic**



# Схема лицензирования

<https://www.couchbase.com/products/editions>

# Вопросы?

# Практика

# ДЗ



# ДЗ

Развернуть кластер Couchbase

Создать БД, наполнить небольшими тестовыми данными

Проверить отказоустойчивость



# Рефлексия

# Рефлексия

Как вам couchbase?

И так как нет доступа в Azure & AWS заменяем на ЯО & SberCloud  
также расшарю предыдущие лекции, если кому актуально

Заполните, пожалуйста,  
опрос о занятии по ссылке в чате  
<https://otus.ru/polls/52770/>

Спасибо за внимание!  
Приходите на следующие вебинары

Аристов Евгений